Kush R. Varshney

IBM Thomas J. Watson Research Center   krvarshn@us.ibm.com
1101 Kitchawan Road   http://krvarshney.github.io
Yorktown Heights, NY 10598   @krvarshney
+1–914–945–1628

Education:

Massachusetts Institute of Technology, Cambridge, MA, USA
Doctor of Philosophy in Electrical Engineering and Computer Science, 2010
Thesis: Frugal Hypothesis Testing and Classification
Committee: Alan S. Willsky (advisor), John W. Fisher, III, Polina Golland, and Joshua B. Tenenbaum
Minor: linguistics

Electrical Engineer, 2010

Master of Science in Electrical Engineering and Computer Science, 2006
Thesis: Joint Anisotropy Characterization and Image Formation in Wide-Angle Synthetic Aperture Radar
Advisors: Müjdat Çetin and John W. Fisher, III
Cumulative GPA 5.00 (A = 5.00)

Cornell University, Ithaca, NY, USA
Bachelor of Science (magna cum laude) in Electrical and Computer Engineering, 2004
Honors Project: Greedy Postprocessing for Spatial Error Concealment in MPEG Video
Advisor: Sheila S. Hemami
Minor: computer science
Cumulative GPA 3.93 (A = 4.00)

Research and Industry Experience:

Distinguished research staff member and manager, *IBM Research*, Thomas J. Watson Research Center, Yorktown Heights, NY, USA, since April 2020.

Principal research staff member and manager, *IBM Research*, Thomas J. Watson Research Center, Yorktown Heights, NY, USA, March 2018–April 2020. (Visiting IBM Research – Africa, Nairobi, Kenya, August 2019–November 2019.)

Research staff member and manager, *IBM Research*, Thomas J. Watson Research Center, Yorktown Heights, NY, USA, December 2016 – March 2018.

Research staff member, *IBM Research*, Thomas J. Watson Research Center, Yorktown Heights, NY, USA, November 2010–December 2016.

Post doctoral researcher, *IBM Research*, Thomas J. Watson Research Center, Yorktown Heights, NY, USA, April–November 2010.

Research assistant, Laboratory for Information and Decision Systems, *Massachusetts Institute of Technology*, Cambridge, MA, USA, October 2004–February 2010.

Intern, Systems and Decision Sciences Section, National Security Engineering Division, *Lawrence Livermore National Laboratory*, Livermore, CA, USA, May–August 2009. (Department of Energy P Clearance.)

Visiting researcher, Laboratoire de Mathématiques Appliquées aux Systèmes, *École Centrale Paris*, Châtenay-Malabry, France, May–August 2006.

Co-op intern, Enterprise Server Products Group, *Sun Microsystems*, Burlington, MA, USA, September 2002–January 2003 and May–August 2003.

Intern, Air Traffic Systems Division, *Sensis Corporation*, DeWitt, NY, USA, May–August 2001.

Summer research volunteer, Nuclear Medicine Division, *State University of New York Health Science Center*, Syracuse, NY, USA, June–September 1998.

Honors:

IBM Research Technical Accomplishments
- Extraordinary Accomplishment
  - Research Contributions to Workforce Innovation & Enterprise Transformation,[†] 2015.
- Outstanding Accomplishment
  - Trustworthy AI,[*] 2019;
  - Computational Creativity, 2014;
  - GMU Proactive Retention, 2013.
- Accomplishment
  - Dictionary- and Sparsity-Driven Imaging,[*] 2018;
  - IBM Leadership in Fight Against Ebola,[*] 2015;
  - WellPoint Health Insurance Exchanges Analytics, 2013;
  - Computational Creativity for Culinary Recipes,[‡] 2013;
  - Analytics-Driven Proactive Retention in the Growth Markets Unit,[†] 2012;
  - Business Impact of Outsourcing Analytics,[†] 2011;
  - Software Group Sales Analytics,[*] 2011.
    [*]Also IBM Outstanding Technical Achievement Award
    [†]Also IBM Research Division Award
    [‡]Also IBM Outstanding Innovation Award

Paper Awards
- Best Paper Addressing Opportunities in AI, Computing Community Consortium / Schmidt Futures Computer Science for Social Good White Paper Competition, 2019.
- Best Research Paper Honorable Mention, SIAM International Conference on Data Mining, 2015.
- Best Social Good Paper Award, ACM SIGKDD Conference on Knowledge Discovery and Data Mining, 2014.
- Best Paper Award, IEEE International Conference on Service Operations and Logistics, and Informatics, 2013.
- Best Student Paper Travel Award, International Conference on Information Fusion, 2009.

Contributor to WellPoint Team, Gerstner Award for Client Excellence (most prestigious internal award at IBM), 2013.

IBM Eminence and Excellence Award
- Research Contributions to Trust and Transparency for AI on the IBM Cloud, 2018.

- WellPoint/IBM Healthcare Predictive Analytics Project, 2012.

National Science Foundation Graduate Research Fellow, 2004–2009.

John McMullen Dean's Scholar, 2000–2004.

Eta Kappa Nu (electrical and computer engineering honor society), inducted 2003.

Tau Beta Pi (engineering honor society), inducted 2003.

Lockheed Martin Award for Academic Excellence (awarded to five outstanding sophomores in School of Electrical and Computer Engineering, Cornell University), 2002.

Rensselaer Medal, 1999.

Research Interests:

Signal processing: detection and estimation theory, sparse signal representation, quantization

Machine learning: interpretable models, fairness and privacy, topological data analysis

Image processing: image formation, image segmentation, level set methods, remote sensing

Applications: workforce analytics, public affairs, olfaction, international development, health care

Publications:

Books and Booklets
AI Fairness: How to Measure and Reduce Unwanted Bias in Machine Learning. Trisha Mahoney, Kush R. Varshney, and Michael Hind. O'Reilly Media, Sebastopol, CA, USA, 2020.

Book Chapters
Learning Interpretable Classification Rules with Boolean Compressed Sensing. Dmitry M. Malioutov, Kush R. Varshney, Amin Emad, and Sanjeeb Dash. *Transparent Data Mining for Big and Small Data*, p. 95–121. Tania Cerquitelli, Daniele Quercia, and Frank Pasquale, editors. Cham, Switzerland: Springer, 2017.

Legislative Prediction with Political and Social Network Analysis. Jun Wang, Kush R. Varshney, and Aleksandra Mojsilović. *Encyclopedia of Social Network Analysis and Mining*, p. 804–811. Reda S. Alhajj and Jon G. Rokne, editors. Heidelberg, Germany: Springer, 2014.

Automatic Fingerprint Matching Systems. Kush R. Varshney. *Glimpses of Systems Theory and Novel Applications: Felicitation Volume in Honour of Professor Raj Kumar Varshney*, p. 149–164. Harjinder Singh Sekhon et al., editors. Aligarh, India: Navin Press, 2005.

Journal and Magazine Articles
FactSheets: Increasing Trust in AI Services through Supplier's Declarations of Conformity. Matthew Arnold, Rachel K. E. Bellamy, Michael Hind, Stephanie Houde, Sameep Mehta, Aleksandra Mojsilović, Ravi Nair, Karthikeyan Natesan Ramamurthy, Alexandra Olteanu, David Piorkowski, Darrell Reimer,

John Richards, Jason Tsay, and Kush R. Varshney. *IBM Journal of Research and Development*, vol. 63, no. 4/5, p. 6, July/September 2019.

AI Fairness 360: An Extensible Toolkit for Detecting and Mitigating Algorithmic Bias. Rachel K. E. Bellamy, Kuntal Dey, Michael Hind, Samuel C. Hoffman, Stephanie Houde, Kalapriya Kannan, Pranay Lohia, Jacquelyn Martino, Sameep Mehta, Aleksandra Mojsilović, Seema Nagar, Karthikeyan Natesan Ramamurthy, John Richards, Diptikalyan Saha, Prasanna Sattigeri, Moninder Singh, Kush R. Varshney, and Yunfeng Zhang. *IBM Journal of Research and Development*, vol. 63, no. 4/5, p. 4, July/September 2019.

Fairness GAN: Generating Datasets with Fairness Properties Using a Generative Adversarial Network. Prasanna Sattigeri, Samuel C. Hoffman, Vijil Chenthamarakshan, and Kush R. Varshney. *IBM Journal of Research and Development*, vol. 63, no. 4/5, p. 3, July/September 2019.

Teaching AI Agents Ethical Values Using Reinforcement Learning and Policy Orchestration. Ritesh Noothigattu, Djallel Bouneffouf, Nicholas Mattei, Rachita Chandra, Piyush Madan, Kush R. Varshney, Murray Campbell, Moninder Singh, and Francesca Rossi. *IBM Journal of Research and Development*, vol. 63, no. 4/5, p. 2, July/September 2019.

Think Your Artificial Intelligence Software is Fair? Think Again. Rachel K. E. Bellamy, Kuntal Dey, Michael Hind, Samuel C. Hoffman, Stephanie Houde, Kalapriya Kannan, Pranay Lohia, Sameep Mehta, Aleksandra Mojsilović, Seema Nagar, Karthikeyan Natesan Ramamurthy, John Richards, Diptikalyan Saha, Prasanna Sattigeri, Moninder Singh, Kush R. Varshney, and Yunfeng Zhang. *IEEE Software*, vol. 36, no. 4, p. 76–80, July–August 2019.

Confronting Data Sparsity to Identify Potential Sources of Zika Virus Spillover Infection Among Primates. Barbara A. Han, Subhabrata Majumdar, Flavio P. Calmon, Benjamin S. Glicksberg, Raya Horesh, Abhishek Kumar, Adam Perer, Elisa B. von Marschall, Dennis Wei, Aleksandra Mojsilović, and Kush R. Varshney. *Epidemics*, vol. 27, p. 59–65, June 2019.

Trustworthy Machine Learning and Artificial Intelligence. Kush R. Varshney. *ACM XRDS Magazine*, vol. 25, no. 3, p. 26–29, Spring 2019.

A Big Data Approach to Computational Creativity: The Curious Case of Chef Watson. Lav R. Varshney, Florian Pinel, Kush R. Varshney, Debarun Bhattacharjya, Angela Schörgendorfer and Yi-Min Chee. *IBM Journal of Research and Development*, vol. 63, no. 1, p. 7, January–February 2019.

Distribution-Preserving k-Anonymity. Dennis Wei, Karthikeyan Natesan Ramamurthy, and Kush R. Varshney. *Statistical Analysis and Data Mining*, vol. 11, no. 6, p. 253–270, December 2018.

Data Pre-Processing for Discrimination Prevention. Flavio P. Calmon, Dennis Wei, Bhanukiran Vinzamuri, Karthikeyan Natesan Ramamurthy, and Kush R. Varshney. *IEEE Journal of Selected Topics in Signal Processing*, vol. 12, no. 5, p. 1106–1119, October 2018.

How to Foster Innovation: A Data-Driven Approach to Measuring Economic Competitiveness. Caitlin Kuhlman, Karthikeyan Natesan Ramamurthy, Prasanna Sattigeri, Aurélie C. Lozano, Lei Cao, Chandra Reddy, Aleksandra Mojsilović, and Kush R. Varshney. *IBM Journal of Research and Development*, vol. 61, no. 6, p. 11, November–December 2017.

Dataflow Representation of Data Analyses: Towards a Platform for Collaborative Data Science. Evan Patterson, Robert McBurney, Hollie Schmidt, Ioana Baldini, Aleksandra Mojsilović, and Kush R. Varshney. *IBM Journal of Research and Development*, vol. 61, no. 6, p. 9, November–December 2017.

Real-Time Understanding of Humanitarian Crises via Targeted Information Retrieval. Kien T. Pham, Prasanna Sattigeri, Amit Dhurandhar, Arpith C. Jacob, Maja Vukovic, Patrice Chataigner, Juliana Freire, Aleksandra Mojsilović, and Kush R. Varshney. *IBM Journal of Research and Development*, vol. 61, no. 6, p. 7, November–December 2017.

Understanding the Ecospace of Philanthropic Projects. Hemank Lamba, Mary E. Helander, Moninder Singh, Nizar Lethif, Anuradha Bhamidipaty, Salman Baset, Aleksandra Mojsilović, and Kush R. Varshney. *IBM Journal of Research and Development*, vol. 61, no. 6, p. 6, November–December 2017.

Effectiveness of Peer Detailing in a Diarrhea Program in Nigeria. Yumeng Tao, Debarun Bhattacharjya, Aliza R. Heching, Aditya Vempaty, Moninder Singh, Felix Lam, Jason Houdek, Mohammed Abubakar, Ahmad Abdulwahab, Tiwadayo Baraimoh, Nnenna Ihebuzor, Aleksandra Mojsilović, and Kush R. Varshney. *IBM Journal of Research and Development*, vol. 61, no. 6, p. 1, November–December 2017.

On the Safety of Machine Learning: Cyber-Physical Systems, Decision Sciences, and Data Products. Kush R. Varshney and Homa Alemzadeh. *Big Data*, vol. 5, no. 3, p. 246–255, September 2017.

Signal Processing for Social Good. Kush R. Varshney. *IEEE Signal Processing Magazine*, vol. 34, no. 3, p. 112, 108, May 2017.

Decision Making with Quantized Priors Leads to Discrimination. Lav R. Varshney and Kush R. Varshney. *Proceedings of the IEEE*, vol. 105, no. 2, p. 241–255, February 2017.

Associative Algorithms for Computational Creativity. Lav R. Varshney, Jun Wang, and Kush R. Varshney. *Journal of Creative Behavior*, vol. 50, no. 3, p. 211–223, September 2016.

Olfactory Signal Processing. Kush R. Varshney and Lav R. Varshney. *Digital Signal Processing*, vol. 48, p. 84–92, January 2016.

Data Challenges in Disease Response: The 2014 Ebola Outbreak and Beyond. Kush R. Varshney, Dennis Wei, Karthikeyan Natesan Ramamurthy, and Aleksandra Mojsilović. *ACM Journal of Data and Information Quality*, vol. 6, no. 2–3, p. 5, June 2015.

Targeting Villages for Rural Development Using Satellite Image Analysis. Kush R. Varshney, George H. Chen, Brian Abelson, Kendall Nowocin, Vivek Sakhrani, Ling Xu, and Brian L. Spatocco. *Big Data*, vol. 3, no. 1, p. 41–53, March 2015.

Optimal Grouping for Group Minimax Hypothesis Testing. Kush R. Varshney and Lav R. Varshney. *IEEE Transactions on Information Theory*, vol. 60, no. 10, p. 6511–6521, October 2014.

Bounded Confidence Opinion Dynamics in a Social Network of Bayesian Decision Makers. Kush R. Varshney. *IEEE Journal of Selected Topics in Signal Processing*, vol. 8, no. 4, p. 576–585, August 2014.

Collaborative Kalman Filtering for Dynamic Matrix Factorization. John Z. Sun, Dhruv Parthasarathy, and Kush R. Varshney. *IEEE Transactions on Signal Processing*, vol. 62, no. 14, p. 3499-3509, July 15, 2014.

Sparsity-Driven Synthetic Aperture Radar Imaging: Reconstruction, Autofocusing, Moving Targets, and Compressed Sensing. Müjdat Çetin, Ivana Stojanović, N. Özben Önhon, Kush R. Varshney, Sadegh Samadi, W. Clem Karl, and Alan S. Willsky. *IEEE Signal Processing Magazine*, vol. 31, no. 4, p. 27–40, July 2014.

Practical Ensemble Classification Error Bounds for Different Operating Points. Kush R. Varshney, Ryan J. Prenger, Tracy L. Marlatt, Barry Y. Chen, and William G. Hanley. *IEEE Transactions on Knowledge and Data Engineering*, vol. 25, no. 11, p. 2590–2601, November 2013.

Sales-Force Performance Analytics and Optimization. Moritz Baier, Jorge E. Carballo, Alice J. Chang, Yingdong Lu, Aleksandra Mojsilović, M. Jonathan Richard, Moninder Singh, Mark S. Squillante, and Kush R. Varshney. *IBM Journal of Research and Development*, vol. 56, no. 6, November/December 2012.

Generalization Error of Linear Discriminant Analysis in Spatially-Correlated Sensor Networks. Kush R. Varshney. *IEEE Transactions on Signal Processing*, vol. 60, no. 6, p. 3295-3301, June 2012.

Bayes Risk Error is a Bregman Divergence. Kush R. Varshney. *IEEE Transactions on Signal Processing*, vol. 59, no. 9, p. 4470–4472, September 2011.

Business Analytics Based on Financial Time Series. Kush R. Varshney and Aleksandra Mojsilović. *IEEE Signal Processing Magazine*, vol. 28, no. 5, p. 83–93, September 2011.

Linear Dimensionality Reduction for Margin-Based Classification: High-Dimensional Data and Sensor Networks. Kush R. Varshney and Alan S. Willsky. *IEEE Transactions on Signal Processing*, vol. 59, no. 6, p. 2496–2512, June 2011.

Classification Using Geometric Level Sets. Kush R. Varshney and Alan S. Willsky. *Journal of Machine Learning Research*, vol. 11, p. 491–516, February 2010.

Postarthroplasty Examination Using X-Ray Images. Kush R. Varshney, Nikos Paragios, Jean-François Deux, Alain Kulski, Rémy Raymond, Phillipe Hernigou, and Alain Rahmouni. *IEEE Transactions on Medical Imaging*, vol. 28, no. 3, p. 469–474, March 2009.

Quantization of Prior Probabilities for Hypothesis Testing. Kush R. Varshney and Lav R. Varshney. *IEEE Transactions on Signal Processing*, vol. 56, no. 10, p. 4553–4562, October 2008.

Sparse Representation in Structured Dictionaries with Application to Synthetic Aperture Radar. Kush R. Varshney, Müjdat Çetin, John W. Fisher, III, and Alan S. Willsky. *IEEE Transactions on Signal Processing*, vol. 56, no. 8, p. 3548–3561, August 2008.

Conference Papers
Inspection of Blackbox Models for Evaluating Vulnerability in Maternal, Newborn, and Child Health. William Ogallo, Skyler Speakman, Victor Akinwande, Kush R. Varshney, Aisha Walcott-Bryant, Charity Wayua, and Komminist Weldemariam. *International Joint Conference on Artificial Intelligence–Pacific Rim International Conference on Artificial Intelligence*, Yokohama, Japan, July 2020. (delayed)

Characterization of Overlap in Observational Studies. Michael Oberst, Fredrik D. Johansson, Dennis Wei, Tian Gao, Gabriel Brat, David Sontag, and Kush R. Varshney. *International Conference on Artificial Intelligence and Statistics*, Palermo, Italy, June 2020. (delayed)

Preservation of Anomalous Subgroups on Variational Autoencoder Transformed Data. Samuel C. Maina, Reginald E. Bryant, William Ogallo, Kush R. Varshney, Skyler Speakman, Celia Cintas, Aisha Walcott-Bryant, Robert-Florian Samoilescu, and Komminist Weldemariam. *IEEE International Conference on Acoustics, Speech, and Signal Processing*, Barcelona, Spain, May 2020. (held virtually)

DADI: Dynamic Discovery of Fair Information with Adversarial Reinforcement Learning. Michiel Bakker, Duy Patrick Tu, Humberto Riverón Valdés, Krishna Gummadi, Kush R. Varshney, Adrian Weller and Alex 'Sandy' Pentland. *ICLR Workshop on Towards Trustworthy ML*, Addis Ababa, Ethiopia, April 2020. (held virtually)

Experiences with Improving the Transparency of AI Models and Services. Michael Hind, Stephanie Houde, Jacquelyn Martino, Aleksandra Mojsilović, David Piorkowski, John Richards, and Kush R. Varshney. *ACM CHI Conference on Human Factors in Computing Systems*, Honolulu, HI, April 2020. (held virtually)

Interpretable Subgroup Discovery in Treatment Effect Estimation with Application to Opioid Prescribing Guidelines. Chirag Nagpal, Dennis Wei, Bhanukiran Vinzamuri, Monica Shekhar, Sara E. Berger, Subhro Das, and Kush R. Varshney. *ACM Conference on Health, Inference, and Learning*, p. 19–29, Toronto, Canada, April 2020. (delayed)

On Mismatched Detection and Safe, Trustworthy Machine Learning. Kush R. Varshney. *Conference on Information Sciences and Systems*, Princeton, NJ, March 2020. (not held, proceedings printed)

Event-Driven Continuous Time Bayesian Networks. Debarun Bhattacharjya, Karthikeyan Shanmugam, Tian Gao, Nicholas Mattei, Kush R. Varshney, and Dharmashankar Subramanian. *AAAI Conference on Artificial Intelligence*, New York, NY, February 2020.

A Natural Language Processing System for Extracting Evidence of Drug Repurposing from Scientific Publications. Shivashankar Subramanian, Ioana Baldini, Sushma Ravichandran, Dmitriy A. Katz-Rogozhnikov, Karthikeyan Natesan Ramamurthy, Prasanna Sattigeri, Kush R. Varshney, Annmarie Wang, Pradeep Mangalath, and Laura B. Kleiman. *Conference on Innovative Applications of Artificial Intelligence*, New York, NY, February 2020.

Data Augmentation for Discrimination Prevention and Bias Disambiguation. Shubham Sharma, Yunfeng Zhang, Jesús M. Ríos Aliaga, Djallel Bouneffouf, Vinod Muthusamy, and Kush R. Varshney. *AAAI/ACM Conference on Artificial Intelligence, Ethics, and Society*, New York, NY, February 2020.

Joint Optimization of AI Fairness and Utility: A Human-Centered Approach. Yunfeng Zhang, Rachel K. E. Bellamy, and Kush R. Varshney. *AAAI/ACM Conference on Artificial Intelligence, Ethics, and Society*, New York, NY, February 2020.

Fair Enough: Improving Fairness in Budget-Constrained Decision Making Using Confidence Thresholds. Michiel Bakker, Duy Patrick Tu, Humberto Riverón Valdés, Krishna Gummadi, Kush R. Varshney, Adrian Weller and Alex 'Sandy' Pentland. *AAAI Workshop on Safe Artificial Intelligence*, New York, NY, February 2020.

How Data Scientists Work Together with Domain Experts in Scientific Collaborations: To Find The Right Answer or to Ask the Right Question? Yaoli Mao, Dakuo Wang, Michael Muller, Kush R. Varshney, Ioana Baldini, Casey Dugan, and Aleksandra Mojsilović. *ACM International Conference on Supporting Group Work*, p. 237, Sanibel Island, FL, January 2020.

Estimating Skin Tone and Effects on Classification Performance in Dermatology Datasets. Newton M. Kinyanjui, Timothy Odonga, Celia Cintas, Noel C. F. Codella, Rameswar Panda, Prasanna Sattigeri, and Kush R. Varshney. *NeurIPS Fair Machine Learning for Health Workshop*, Vancouver, Canada, December 2019.

Drug Repurposing for Cancer: An NLP Approach to Identify Low-Cost Therapies. Shivashankar Subramanian, Ioana Baldini, Sushma Ravichandran, Dmitriy A. Katz-Rogozhnikov, Karthikeyan Natesan Ramamurthy, Prasanna Sattigeri, Kush R. Varshney, Annmarie Wang, Pradeep Mangalath, and Laura B. Kleiman. *NeurIPS Workshop on Machine Learning for Health*, Vancouver, Canada, December 2019.

DADI: Dynamic Discovery of Fair Information with Adversarial Reinforcement Learning. Michiel Bakker, Duy Patrick Tu, Humberto Riverón Valdés, Krishna Gummadi, Kush R. Varshney, Adrian Weller and Alex 'Sandy' Pentland. *NeurIPS Workshop on Human-Centric Machine Learning*, Vancouver, Canada, December 2019.

Subgroup Preservation in Financial Data Anonymized by a Variational Autoencoder. Samuel C. Maina, Reginald E. Bryant, William Ogallo, Kush R. Varshney, Skyler Speakman, Celia Cintas, Aisha Walcott-Bryant, and Robert-Florian Samoilescu. *NeurIPS Workshop on Robust AI in Financial Services: Data, Fairness, Explainability, Trustworthiness, and Privacy*, Vancouver, Canada, December 2019.

Teaching AI Ethical Values Using Reinforcement Learning and Policy Orchestration. Ritesh Noothigattu, Djallel Bouneffouf, Nicholas Mattei, Rachita Chandra, Piyush Madan, Kush R. Varshney, Murray Campbell, Moninder Singh, and Francesca Rossi. *International Joint Conference on Artificial Intelligence*, p. 6377–6381, Macao, August 2019.

Event-Driven Continuous Time Bayesian Networks: An Application in Modeling Progression out of Poverty through Integrated Social Services. Debarun Bhattacharjya, Karthikeyan Shanmugam, Tian Gao, Nicholas Mattei, and Kush R. Varshney. *IJCAI Workshop on AI for Social Good*, Macau, August 2019.

On Fairness in Budget-Constrained Decision Making. Michiel Bakker, Alejandro Noriega Campero, Duy Patrick Tu, Prasanna Sattigeri, Kush R. Varshney, and Alex 'Sandy' Pentland. *KDD Workshop on Explainable Artificial Intelligence*, Anchorage, AK, August 2019.

Open Platforms for Artificial Intelligence for Social Good: Common Patterns as a Pathway to True Impact. Kush R. Varshney and Aleksandra Mojsilović. *ICML Workshop on AI for Social Good*, Long Beach, CA, June 2019.

Teaching AI to Explain its Decisions Using Embeddings and Multi-Task Learning. Noel C. F. Codella, Michael Hind, Karthikeyan Natesan Ramamurthy, Murray Campbell, Amit Dhurandhar, Kush R. Varshney, Dennis Wei, and Aleksandra Mojsilović. *ICML Workshop on Human in the Loop Learning*, Long Beach, CA, June 2019.

Topological Data Analysis of Decision Boundaries with Application to Model Selection. Karthikeyan Natesan Ramamurthy, Kush R. Varshney, and Krishnan Mody. *International Conference on Machine Learning*, Long Beach, CA, p. 5351–5360, June 2019.

Bias Mitigation Post-Processing for Individual and Group Fairness. Pranay K. Lohia, Karthikeyan Natesan Ramamurthy, Manish Bhide, Diptikalyan Saha, Kush R. Varshney, and Ruchir Puri. *IEEE International Conference on Acoustics, Speech and Signal Processing*, Brighton, UK, p. 2847–2851, May 2019.

Constructing and Compressing Frames in Blockchain-Based Verifiable Multi-Party Computation. Ravi Kiran Raman, Kush R. Varshney, Roman Vaculin, Nelson Kibichii Bore, Sekou L. Remy, Eleftheria K. Pissadaki, and Michael Hind. *IEEE International Conference on Acoustics, Speech and Signal Processing*, Brighton, UK, p. 7500–7504, May 2019.

Promoting Distributed Trust in Machine Learning and Computational Simulation. Nelson Kibichii Bore, Ravi Kiran Raman, Isaac M. Markus, Sekou L. Remy, Oliver Bent, Michael Hind, Eleftheria K. Pissadaki, Biplav Srivastava, Roman Vaculin, Kush R. Varshney, and Komminist Weldemariam. *IEEE International Conference on Blockchain and Cryptocurrency*, Seoul, Korea, May 2019.

A Scalabale Blockchain Approach for Trusted Computation and Verifiable Simulation in Multi-Party Collaboration. Ravi Kiran Raman, Roman Vaculin, Michael Hind, Sekou L. Remy, Eleftheria K. Pissadaki, Nelson Kibichii Bore, Roozbeh Daneshvar, Biplav Srivastava, and Kush R. Varshney. *IEEE International Conference on Blockchain and Cryptocurrency*, Seoul, Korea, May 2019.

Fairness GAN: Generating Datasets with Fairness Properties using a Generative Adversarial Network. Prasanna Sattigeri, Samuel C. Hoffman, Vijil Chenthamarakshan, and Kush R. Varshney. *ICLR Workshop on Safe Machine Learning*, New Orleans, LA, May 2019.

Fair Transfer Learning with Missing Protected Attributes. Amanda Coston, Karthikeyan Natesan Ramamurthy, Dennis Wei, Kush R. Varshney, Skyler Speakman, Zairah Mustahsan, and Supriyo Chakraborty. *AAAI/ACM Conference on Artificial Intelligence, Ethics, and Society*, Honolulu, HI, p. 91–98, January 2019.

TED: Teaching AI to Explain Its Decisions. Michael Hind, Dennis Wei, Murray Campbell, Noel C. F. Codella, Amit Dhurandhar, Aleksandra Mojsilović, Karthikeyan Natesan Ramamurthy, and Kush R. Varshney. *AAAI/ACM Conference on Artificial Intelligence, Ethics, and Society*, Honolulu, HI, p. 123–129, January 2019.

Financial Forecasting and Analysis for Low-Wage Workers. Wenyu Zhang, Raya Horesh, Karthikeyan Natesan Ramamurthy, Lingfei Wu, Jinfeng Yi, Kryn Anderson, and Kush R. Varshney. *Data for Good Exchange*, New York, NY, September 2018.

Teaching Machines to Understand Data Science Code by Semantic Enrichment of Dataflow Graphs. Evan Patterson, Ioana Baldini, Aleksandra Mojsilović, and Kush R. Varshney. *KDD Workshop on Fragile Earth: Theory Guided Data Science to Enhance Scientific Discovery*, London, UK, p. 5, August 2018.

Semantic Representation of Data Science Programs. Evan Patterson, Ioana Baldini, Aleksandra Mojsilović, and Kush R. Varshney. *International Joint Conference on Artificial Intelligence and European Conference on Artificial Intelligence*, Stockholm, Sweden, p. 5847–5849, July 2018.

Why Interpretability in Machine Learning? An Answer Using Distributed Detection and Data Fusion Theory. Kush R. Varshney, Prashant Khanduri, Pranay Sharma, Shan Zhang, and Pramod K. Varshney. *ICML Workshop on Human Interpretability in Machine Learning*, p. 15–20, Stockholm, Sweden, July 2018.

The Effect of Extremist Violence on Hateful Speech Online. Alexandra Olteanu, Carlos Castillo, Jeremy Boy, and Kush R. Varshney. *International AAAI Conference on Weblogs and Social Media*, Palo Alto, CA, p. 221–230, June 2018.

False Positive Control with Concave Penalties using Stability Selection. Bhanukiran Vinzamuri and Kush R. Varshney. *IEEE Data Science Workshop*, Lausanne, Switzerland, p. 76–80, June 2018.

Assessing National Development Plans for Alignment with Sustainable Development Goals via Semantic Search. Jonathan Galsurkar, Moninder Singh, Lingfei Wu, Aditya Vempaty, Mikhail Sushkov, Devika Iyer, Serge Kapto, and Kush R. Varshney. *Conference on Innovative Applications of Artificial Intelligence*, New Orleans, LA, p. 7753–7758, February 2018.

Neurology-as-a-Service for the Developing World. Tejas Dharamsi, Payel Das, Tejaswini Pedapati, Gregory Bramble, Vinod Muthusamy, Horst Samulowitz, and Kush R. Varshney, Yuvaraj Rajamanickam, John Thomas, and Justin Dauwels. *NIPS Workshop on Machine Learning for the Developing World*, Long Beach, CA, December 2017.

Scalable Demand-Aware Recommendation. Jinfeng Yi, Cho-Jui Hsieh, Kush R. Varshney, Lijun Zhang, and Yao Li. *Advances in Neural Information Processing Systems*, Long Beach, CA, p. 2412–2421, December 2017.

Optimized Pre-Processing for Discrimination Prevention. Flavio P. Calmon, Dennis Wei, Karthikeyan Natesan Ramamurthy, Bhanukiran Vinzamuri, and Kush R. Varshney. *Advances in Neural Information Processing Systems*, Long Beach, CA, p. 3992–4001, December 2017.

Exploring the Causal Relationships between Initial Opioid Prescriptions and Outcomes. Jinghe Zhang, Vijay S. Iyengar, Dennis Wei, Bhanukiran Vinzamuri, Hamsa Bastani, Alexander R. Macalalad, Anne E. Fischer, Gigi Yuen-Reed, Aleksandra Mojsilović, and Kush R. Varshney. *AMIA Workshop on Data Mining for Medical Informatics*, Washington, DC, November 2017.

An End-To-End Machine Learning Pipeline That Ensures Fairness Policies. Samiulla Shaikh, Harit Vishwakarma, Sameep Mehta, Kush R. Varshney, Karthikeyan Natesan Ramamurthy, and Dennis Wei. *Data for Good Exchange*, New York, NY, September 2017.

The Limits of Abstract Evaluation Metrics: The Case of Hate Speech Detection. Alexandra Olteanu, Kartik Talamadupula, and Kush R. Varshney. *ACM Web Science Conference*, Troy, NY, p. 405–406, June 2017.

Statistical Analysis of Peer Detailing for Children's Diarrhea Treatments. Yumeng Tao, Debarun Bhattacharjya, Aliza R. Heching, Aditya Vempaty, Moninder Singh, Felix Lam, Kush R. Varshney, and Aleksandra Mojsilović. *AAAI Spring Symposium on AI for Social Good*, Stanford, CA, p. 101–106, March 2017.

Machine Representation of Data Analyses: Towards a Platform for Collaborative Data Science. Evan J. Patterson, Ioana Baldini, Aleksandra Mojsilović, and Kush R. Varshney. *AAAI Spring Symposium on AI for Social Good*, Stanford, CA, p. 53–59, March 2017.

Information Retrieval, Fusion, Completion, and Clustering for Employee Expertise Estimation. Raya Horesh, Kush R. Varshney, and Jinfeng Yi. *IEEE International Conference on Big Data*, Washington, DC, p. 1385–1393, December 2016.

Stable Estimation of Granger-Causal Factors of Country-Level Innovation. Aurélie C. Lozano, Prasanna Sattigeri, Aleksandra Mojsilović, and Kush R. Varshney. *IEEE Global Conference on Signal and Information Processing*, Washington, DC, p. 1290–1294, December 2016.

Interpretable Machine Learning via Convex Cardinal Shape Composition. Kush R. Varshney. *Allerton Conference on Communication, Control, and Computing*, Monticello, IL, p. 327–330, September 2016.

Learning Sparse Two-Level Boolean Rules. Guolong Su, Dennis Wei, Kush R. Varshney, and Dmitry M. Malioutov. *IEEE Workshop on Machine Learning for Signal Processing*, Salerno, Italy, September 2016.

Dynamic Matrix Factorization with Social Influence. Aleksandr Y. Aravkin, Kush R. Varshney, and Liu Yang. *IEEE Workshop on Machine Learning for Signal Processing*, Salerno, Italy, September 2016.

Understanding Innovation to Drive Sustainable Development. Prasanna Sattigeri, Aurélie Lozano, Aleksandra Mojsilović, Kush R. Varshney, and Mahmoud Naghshineh. *ICML Workshop on #Data4Good: Machine Learning in Social Good Applications*. New York, NY, p. 21–25, June 2016.

Interpretable Two-Level Boolean Rule Learning for Classification. Guolong Su, Dennis Wei, Kush R. Varshney, and Dmitry M. Malioutov. *ICML Workshop on Human Interpretability in Machine Learning*, New York, NY, p. 66–70, June 2016.

Fidelity Loss in Distribution-Preserving Anonymization and Histogram Equalization. Lav R. Varshney and Kush R. Varshney. *Conference on Information Sciences and Systems*, Princeton, NJ, p. 30–35, March 2016.

Engineering Safety in Machine Learning. Kush R. Varshney. *Information Theory and Applications Workshop*, La Jolla, CA, February 2016.

Predictive Modeling of Customer Repayment for Sustainable Pay-As-You-Go Solar Power in Rural India. Hugo Gerard, Kamalesh Rao, Mark Simithraaratchy, Kush R. Varshney, Kunal Kabra, and G. Paul Needham. *Data for Good Exchange Conference*, New York, NY, September 2015.

Data Science of the People, for the People, by the People: A Viewpoint on an Emerging Dichotomy. Kush R. Varshney. *Data for Good Exchange Conference*, New York, NY, September 2015.

From Open Data Ecosystems to Systems of Innovation: A Journey to Realize the Promise of Open Data. Shubir Kapoor, Aleksandra Mojsilović, Jade Nguyen Strattner, and Kush R. Varshney. *Data for Good Exchange Conference*, New York, NY, September 2015.

A Robust Nonlinear Kalman Smoothing Approach for Dynamic Matrix Factorization. Aleksandr Y. Aravkin, Kush R. Varshney, and Dmitry M. Malioutov. *Signal Processing with Adaptive Sparse Structured Representations Workshop*, Cambridge, United Kingdom, July 2015.

A Semiquantitative Group Testing Approach for Learning Interpretable Clinical Prediction Rules. Amin Emad, Kush R. Varshney, and Dmitry M. Malioutov. *Signal Processing with Adaptive Sparse Structured Representations Workshop*, Cambridge, United Kingdom, July 2015.

Optigrow: People Analytics for Job Transfers. Dennis Wei, Kush R. Varshney, and Marcy Wagman. *IEEE International Congress on Big Data*, New York, NY, p. 535–542, June–July 2015.

Health Insurance Market Risk Assessment: Covariate Shift and k-Anonymity. Dennis Wei, Karthikeyan Natesan Ramamurthy, and Kush R. Varshney. *SIAM International Conference on Data Mining*, Vancouver, Canada, p. 226–234, April–May 2015. (*Best Research Paper Honorable Mention*)

Learning Interpretable Classification Rules Using Sequential Row Sampling. Sanjeeb Dash, Dmitry M. Malioutov, and Kush R. Varshney. *IEEE International Conference on Acoustics, Speech and Signal Processing*, Brisbane, Australia, p. 3337–3341, April 2015.

Robust Binary Hypothesis Testing Under Contaminated Likelihoods. Dennis Wei and Kush R. Varshney. *IEEE International Conference on Acoustics, Speech and Signal Processing*, Brisbane, Australia, p. 3407–3411, April 2015.

Persistent Topology of Decision Boundaries. Kush R. Varshney and Karthikeyan Natesan Ramamurthy. *IEEE International Conference on Acoustics, Speech and Signal Processing*, Brisbane, Australia, p. 3931–3935, April 2015.

Targeting Direct Cash Transfers to the Extremely Poor. Brian Abelson, Kush R. Varshney, and Joy Sun. *ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, New York, NY, p. 1563–1572, August 2014. *(Best Social Good Paper Award)*

Predicting Employee Expertise for Talent Management in the Enterprise. Kush R. Varshney, Vijil Chenthamarakshan, Scott W. Fancher, Jun Wang, Dongping Fang, and Aleksandra Mojsilović. *ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, New York, NY, p. 1729–1738, August 2014.

An Analysis of Losing Unimportant Points in Tennis. Kush R. Varshney. *KDD Workshop on Large-Scale Sports Analytics*, New York, NY, August 2014.

Computing Persistent Homology Under Random Projection. Karthikeyan Natesan Ramamurthy, Kush R. Varshney, and Jayaraman J. Thiagarajan. *IEEE International Workshop on Statistical Signal Processing*, Gold Coast, Australia, p. 105–108, June–July 2014.

Food Steganography with Olfactory White. Kush R. Varshney and Lav R. Varshney. *IEEE International Workshop on Statistical Signal Processing*, Gold Coast, Australia, p. 21–24, June–July 2014.

Active Odor Cancellation. Kush R. Varshney and Lav R. Varshney. *IEEE International Workshop on Statistical Signal Processing*, Gold Coast, Australia, p. 25–28, June–July 2014.

Screening for Learning Classification Rules via Boolean Compressed Sensing. Sanjeeb Dash, Dmitry M. Malioutov, and Kush R. Varshney. *IEEE International Conference on Acoustics, Speech, and Signal Processing*, Florence, Italy, p. 3360–3364, May 2014.

Prescriptive Analytics for Allocating Sales Teams to Opportunities. Ban Kawas, Mark S. Squillante, Dharmashankar Subramanian, and Kush R. Varshney. *IEEE International Conference on Data Mining Workshops*, Dallas, TX, p. 211–218, December 2013.

Quantifying and Recommending Expertise When New Skills Emerge. Dongping Fang, Kush R. Varshney, Jun Wang, Karthikeyan Natesan Ramamurthy, Aleksandra Mojsilović, and John H. Bauer. *IEEE International Conference on Data Mining Workshops*, Dallas, TX, p. 672–679, December 2013.

Quantile Regression for Workforce Analytics. Karthikeyan Natesan Ramamurthy, Kush R. Varshney, and Moninder Singh. *IEEE Global Conference on Signal and Information Processing*, Austin, TX, p. 1134, December 2013.

Flavor Pairing in Medieval European Cuisine: A Study in Cooking with Dirty Data. Kush R. Varshney, Lav R. Varshney, Jun Wang, and Daniel Meyers. *IJCAI Workshop on Cooking with Computers*, Beijing, China, p. 3–12, August 2013.

Balancing Lifetime and Classification Accuracy of Wireless Sensor Networks. Kush R. Varshney and Peter M. van de Ven. *ACM International Symposium on Mobile Ad Hoc Networking and Computing*, Bengaluru, India, p. 31–38, July–August 2013.

Expertise Assessment with Multi-Cue Semantic Information. Jun Wang, Kush R. Varshney, Aleksandra Mojsilović, Dongping Fang, and John H. Bauer. *IEEE International Conference on Service Operations and Logistics, and Informatics*, Dongguan, China, p. 534–539, July 2013. *(Best Paper Award)*

Dose-Response Signal Estimation and Optimization for Salesforce Management. Kush R. Varshney and Moninder Singh. *IEEE International Conference on Service Operations and Logistics, and Informatics*, Dongguan, China, p. 328–333, July 2013.

Cognition as a Part of Computational Creativity. Lav R. Varshney, Florian Pinel, Kush R. Varshney, Angela Schörgendorfer, and Yi-Min Chee. *IEEE International Conference on Cognitive Informatics and Cognitive Computing*, New York, NY, p. 36–43, July 2013.

Predicting and Recommending Skills in the Social Enterprise. Kush R. Varshney, Jun Wang, Aleksandra Mojsilović, Dongping Fang, and John H. Bauer. *International AAAI Conference on Weblogs and Social Media Workshop on Social Computing for Workforce 2.0*, Cambridge, MA, p. 20–23, July 2013.

A Salesforce Control Theory Analysis of Enterprise Microblog Posts. Kush R. Varshney and N. Sadat Shami. *International AAAI Conference on Weblogs and Social Media Workshop on Social Computing for Workforce 2.0*, Cambridge, MA, p. 16–19, July 2013.

Exact Rule Learning via Boolean Compressed Sensing. Dmitry M. Malioutov and Kush R. Varshney. *International Conference on Machine Learning*, Atlanta, GA, p. 765–773, June 2013.

Opinion Dynamics with Bounded Confidence in the Bayes Risk Error Divergence Sense. Kush R. Varshney. *IEEE International Conference on Acoustics, Speech, and Signal Processing*, Vancouver, Canada, p. 6600–6604, May 2013.

Interactive Visual Salesforce Analytics. Kush R. Varshney, Jamie C. Rasmussen, Aleksandra Mojsilović, Moninder Singh, and Joan M. DiMicco. *International Conference on Information Systems*, Orlando, FL, December 2012.

An Analytics Approach for Proactively Combating Voluntary Attrition of Employees. Moninder Singh, Kush R. Varshney, Jun Wang, Aleksandra Mojsilović, Alisia R. Gill, Patricia I. Faur, and Raphael Ezry. *IEEE International Conference on Data Mining Workshops*, Brussels, Belgium, p. 317–323, December 2012.

Deconvolving the Productivity of Salespeople via Constrained Quadratic Programming. Gautam K. Bhat and Kush R. Varshney. *National Systems Conference*, Annamalainagar, India, December 2012.

Decision Trees for Heterogeneous Dose-Response Signal Analysis. Kush R. Varshney, Moninder Singh, and Jun Wang. *IEEE International Workshop on Statistical Signal Processing*, Ann Arbor, MI, p. 916–919, August 2012.

Does Selection Bias Blind Performance Diagnostics of Business Decision Models? A Case Study in Salesforce Optimization. Jun Wang, Moninder Singh, and Kush R. Varshney. *IEEE International Conference on Service Operations and Logistics, and Informatics*. Suzhou, China, p. 416–421, July 2012.

Legislative Prediction via Random Walks over a Heterogeneous Graph. Jun Wang, Kush R. Varshney, and Aleksandra Mojsilović. *SIAM International Conference on Data Mining*, Anaheim, CA, p. 1095–1106, April 2012.

Dynamic Matrix Factorization: A State Space Approach. John Z. Sun, Kush R. Varshney, and Karthik Subbian. *IEEE International Conference on Acoustics, Speech, and Signal Processing*, Kyoto, Japan, p. 1897–1900, March 2012.

Estimating Post-Event Seller Productivity Profiles in Dynamic Sales Organizations. Kush R. Varshney, Moninder Singh, Mayank Sharma, and Aleksandra Mojsilović. *IEEE International Conference on Data Mining Workshops*, Vancouver, Canada, p. 1191–1198, December 2011.

A Risk Bound for Ensemble Classification with a Reject Option. Kush R. Varshney. *IEEE International Workshop on Statistical Signal Processing*, Nice, France, p. 773–776, June 2011.

Multilevel Minimax Hypothesis Testing. Kush R. Varshney and Lav R. Varshney. *IEEE International Workshop on Statistical Signal Processing*, Nice, France, p.109–112, June 2011.

MCMC Inference of the Shape and Variability of Time-Response Signals. Dmitriy A. Katz-Rogozhnikov, Kush R. Varshney, Aleksandra Mojsilović, and Moninder Singh. *IEEE International Conference on Acoustics, Speech, and Signal Processing*, p. 3956–3959, Prague, Czech Republic, May 2011.

Spatially-Correlated Sensor Discriminant Analysis. Kush R. Varshney. *IEEE International Conference on Acoustics, Speech, and Signal Processing*, p. 3680–3683, Prague, Czech Republic, May 2011.

Categorical Decision Making by People, Committees, and Crowds. Lav R. Varshney, Joong Bum Rhim, Kush R. Varshney, and Vivek K Goyal. *Information Theory and Applications Workshop*, La Jolla, CA, February 2011.

Class-Specific Error Bounds for Ensemble Classifiers. Ryan J. Prenger, Tracy D. Lemmond, Kush R. Varshney, Barry Y. Chen, and William G. Hanley. *ACM SIGKDD Conference on Knowledge Discovery and Data Mining*, p. 843–852, Arlington, VA, July 2010.

Learning Dimensionality-Reduced Classifiers for Information Fusion. Kush R. Varshney and Alan S. Willsky. *International Conference on Information Fusion*, p. 1881–1888, Seattle, WA, July 2009. *(Best Student Paper Travel Award)*

Supervised Learning of Classifiers via Level Set Segmentation. Kush R. Varshney and Alan S. Willsky. *IEEE Workshop on Machine Learning for Signal Processing*, p. 115–120, Cancún, Mexico, October 2008.

Minimum Mean Bayes Risk Error Quantization of Prior Probabilities. Kush R. Varshney and Lav R. Varshney. *IEEE International Conference on Acoustics, Speech, and Signal Processing*, p. 3445–3448, Las Vegas, NV, April 2008.

Geniş Açılı Radarda Görüntü Oluşturma ve Yönbağımlılık Tespiti için Seyrek Sinyal Temsiline Dayalı bir Yaklaşım (A Sparse Signal Representation-Based Approach to Image Formation and Anisotropy Determination in Wide-Angle Radar). Kush R. Varshney, Müjdat Çetin, John W. Fisher, III, and Alan S. Willsky. *IEEE Signal Processing and Communication Applications Conference*, Eskişehir, Turkey, June 2007.

Multi-View Stereo Reconstruction of Total Knee Replacement from X-Rays. Kush R. Varshney, Nikos Paragios, Alain Kulski, Remy Raymond, Phillipe Hernigou, and Alain Rahmouni. *IEEE International Symposium on Biomedical Imaging: From Nano to Macro*, p. 1148–1151, Arlington, VA, April 2007.

Wide-Angle SAR Image Formation with Migratory Scattering Centers and Regularization in Hough Space. Kush R. Varshney, Müjdat Çetin, John W. Fisher, III, and Alan S. Willsky. *Annual Workshop on Adaptive Sensor Array Processing*, Lexington, MA, June 2006.

Joint Image Formation and Anisotropy Characterization in Wide-Angle SAR. Kush R. Varshney, Müjdat Çetin, John W. Fisher, III, and Alan S. Willsky. *SPIE Defense and Security Symposium, Algorithms for Synthetic Aperture Radar Imagery XIII*, p. 62370D, Orlando (Kissimmee), FL, April 2006.

Conference Presentations and Posters
Empirical and Theoretical Relationships Among Pillars of Trustworthy Machine Learning. Kush R. Varshney and Moninder Singh. *INFORMS Annual Meeting*, National Harbor, MD, November 2020.

Drug Repurposing for Cancer: An NLP Approach to Identify Low-Cost Therapies. Shivashankar Subramanian, Ioana Baldini, Sushma Ravichandran, Dmitriy A. Katz-Rogozhnikov, Karthikeyan Natesan Ramamurthy, Prasanna Sattigeri, Kush R. Varshney, Annmarie Wang, Pradeep Mangalath, and Laura B. Kleiman. *New York Academy of Sciences Natural Language, Dialog and Speech Symposium*, New York, NY, November 2019.

Rigorous Analysis of Racial Bias in Gender Classification. Vidya Muthukumar, Tejaswini Pedapati, Nalini Ratha, Prasanna Sattigeri, Chai-Wah Wu, Brian Kingsbury, Abhishek Kumar, Samuel Thomas, Aleksandra Mojsilović, and Kush R. Varshney. *Women in Machine Learning Workshop*, Montreal, Canada, December 2018.

SimplerVoice: A Key Message & Visual Description Generator System for Illiteracy. Minh N. B. Nguyen, Samuel Thomas, Kush R. Varshney, Sujatha Kashyap, and Anne E. Gattiker. *Data for Good Exchange*, New York, NY, September 2018.

Decision Support for Policymaking: Causal Inference Algorithm and Case Study. Bhanukiran Vinzamuri, Aleksandra Mojsilović, and Kush R. Varshney. *Workshop on Mechanism Design for Social Good*, Ithaca, NY, June 2018.

How to Foster Innovation: A Data-Driven Approach to Measuring Economic Competitiveness. Caitlin Kuhlman, Karthikeyan Natesan Ramamurthy, Prasanna Sattigeri, Aurélie C. Lozano, Lei Cao, Chandra Reddy, Aleksandra Mojsilović, and Kush R. Varshney. *Workshop on Mechanism Design for Social Good*, Ithaca, NY, June 2018.

Decision Making With Quantized Priors Leads to Discrimination. Lav R. Varshney and Kush R. Varshney. *Illinois Summit on Diversity in Psychological Science*, Champaign, IL, March, 2018.

A Social Good Program at an Industrial Research Laboratory. Kush R. Varshney. *American Association for the Advancement of Science Annual Meeting*, Austin, TX, February, 2018.

SimplerVoice: A Key Message & Visual Description Generator System for Illiteracy. Minh N. B. Nguyen, Samuel Thomas, Anne E. Gattiker, Sujatha Kashyap, and Kush R. Varshney. *Women in Machine Learning Workshop*, Long Beach, CA, December 2017.

Machine Learning from Health Insurance Administrative Data: Opioids, Obamacare, and Other Applications. Kush R. Varshney. *Administrative Data Research Facilities Network Conference*, Washington, DC, November 2017.

Interpretable Two-Level Boolean Rule Learning. Guolong Su, Dennis Wei, Kush R. Varshney, and Dmitry M. Malioutov. *INFORMS Annual Meeting*, Houston, TX, October 2017.

Semantic Searching for Efficient Assessment of Sustainable Development in National Plans. Jonathan Galsurkar, Aditya Vempaty, Kush R. Varshney, Lingfei Wu, Mikhail Sushkov, Moninder Singh, Devika Iyer, and Serge Kapto. *Data Science for Social Good Conference*, Chicago, IL, September 2017.

Cognitive Disease Hunter: Developing Automated Pathogen Feature Extraction from Scientific Literature. Timothy NeCamp, Prasanna Sattigeri, Dennis Wei, Emily Ray, Youssef Drissi, Ananya Poddar, Diwakar Mahajan, Sarah Bowden, Barbara A. Han, Aleksandra Mojsilović, and Kush R. Varshney. *Data Science for Social Good Conference*, Chicago, IL, September 2017.

Demystifying Social Entrepreneurship: An NLP Based Approach to Finding a Social Good Fellow. Aditya Garg, Alexandra Olteanu, Richard B. Segal, Dmitriy A. Katz-Rogozhnikov, Keerthana Kumar, Joana Maria, Liza Mueller, Ben Beers, and Kush R. Varshney. *Data Science for Social Good Conference*, Chicago, IL, September 2017.

Decision Support for Policymaking: Causal Inference Algorithm and Case Study. Bhanukiran Vinzamuri, Aleksandra Mojsilović, and Kush R. Varshney. *Data Science for Social Good Conference*, Chicago, IL, September 2017.

A Bayesian Approach for Predicting Neotropical Primate Reservoirs of Zika Virus. Subhabrata Majumdar, Dennis Wei, Adam Perer, Benjamin Glicksberg, Aleksandra Mojsilović, Kush R. Varshney, and Barabara A. Han. *International Symposium on Business and Industrial Statistics*, Yorktown Heights, NY, June 2017.

Employee Expertise Estimation by Information Retrieval, Fusion, Completion, and Clustering. Raya Horesh, Kush R. Varshney, and Jinfeng Yi. *SDM Workshop on Women in Data Science*, Houston, TX, April 2017.

Persistent Homology of Classifier Decision Boundaries. Kush R. Varshney and Karthikeyan Natesan Ramamurthy. *STOC/SoCG Workshop on Geometry and Machine Learning*, Cambridge, MA, June 2016.

Learning Interpretable Clinical Prediction Rules using Threshold Group Testing. Amin Emad, Kush R. Varshney, and Dmitry M. Malioutov. *NIPS Workshop on Machine Learning in Healthcare*, Montréal, Canada, December 2015.

Assessing Expertise in the Enterprise: The Recommender Point of View. Aleksandra Mojsilović and Kush R. Varshney. *ACM Recommender Systems Conference*, p. 231, Vienna, Austria, September 2015.

Learning Interpretable Classification Rules via Boolean Compressed Sensing. Dmitry M. Malioutov, Kush R. Varshney, and Sanjeeb Dash. *Joint Statistical Meetings*, Seattle, WA, August 2015.

Personalization of Product Novelty Assessment via Bayesian Surprise. Nan Shao, Kush R. Varshney, Lav R. Varshney, Florian Pinel, Anshul Sheopuri, and Pavankumar Murali. *Joint Statistical Meetings*, Boston, MA, August 2014.

Talent Analytics to Predict Employee Job Roles and Skill Sets Using Diverse Data Sources. Kush R. Varshney, Vijil Chenthamarakshan, Jun Wang, Dongping Fang, and Aleksandra Mojsilović. *International Symposium on Business and Industrial Statistics / Conference of the ASA Section on Statistical Learning and Data Mining*, Durham, NC, June 2014.

Sales Team Effort Allocation Analytics. Ban Kawas, Mark S. Squillante, Dharmashankar Subramanian, and Kush R. Varshney. *INFORMS Annual Meeting*, Minneapolis, MN, October 2013.

Dynamic Factor Modeling via Robust Subspace Tracking. Aleksandr Y. Aravkin, Kush R. Varshney, and Dmitry M. Malioutov. *Industrial-Academic Workshop on Optimization in Finance and Risk Management*, Toronto, Canada, September 2013.

More Contentious Issues Lead to More Factions: Bounded Confidence Opinion Dynamics of Bayesian Decision Makers. Kush R. Varshney. *Interdisciplinary Workshop on Information and Decision in Social Networks*, Cambridge, MA, November 2012.

Proactive Retention. Yingdong Lu, Aleksandra Mojsilović, Moninder Singh, Mark S. Squillante, Kush R. Varshney, and Jun Wang. *INFORMS Annual Meeting*, Phoenix, AZ, October 2012.

Heterogeneous Graphs in Social Business. Kush R. Varshney, Jun Wang, and Aleksandra Mojsilović. *Graph Exploitation Symposium*, Dedham, MA, April, 2012.

Map of the Marketplace: Visualizing the Relationships in the IT Services Marketplace. Aleksandra Mojsilović, Kush R. Varshney, and Jun Wang. *INFORMS Annual Meeting*, Charlotte, NC, November 2011.

Identifying Optimal Sales Team Composition for Business Opportunities. Aleksandra Mojsilović, Moninder Singh, Kush R. Varshney, and Jun Wang. *INFORMS Annual Meeting*, Charlotte, NC, November 2011.

Classification of IT Service Tickets for Defect Prevention. Kush R. Varshney, Dongping Fang, Aliza R. Heching, Aleksandra Mojsilović, and Moninder Singh. *INFORMS Annual Meeting*, Charlotte, NC, November 2011.

A Framework for Sales Force Productivity Profile Estimation. Mayank Sharma, Aleksandra Mojsilović, Moninder Singh, and Kush R. Varshney. *INFORMS Annual Meeting*, Austin, TX, November 2010.

Random Forest Classifier Performance in Low False Alarm and Low Missed Detection Regimes. Kush R. Varshney. *Lawrence Livermore National Laboratory Institutional Summer Student Poster Symposium*, Livermore, CA, August 2009.

Block-Segmentation and Classification of Grayscale Postal Images. Kush R. Varshney. *Bits on our Mind Symposium*, Ithaca, NY, March 2004.

Theses and Technical Reports

Frugal Hypothesis Testing and Classification. Kush R. Varshney. *Ph.D. Thesis*. Massachusetts Institute of Technology, Cambridge, MA, February 2010.

The Feature Analysis of Flow-Based Statistics for Network Traffic Classification. Kelley B. Herndon Ford, Kush R. Varshney, Tracy D. Lemmond, William G. Hanley, Barry Y. Chen, and William C. Kallander. *Technical Report*. Lawrence Livermore National Laboratory, Livermore, CA, September 2009.

Surface Evolution for 3-D Shape Reconstruction of Knee Joint Prosthetics and Bones. Kush R. Varshney and Nikos Paragios. *Technical Report 0703*. Laboratoire de Mathématiques Appliquées aux Systèmes, École Centrale Paris, Châtenay-Malabry, France, February 2007.

Joint Anisotropy Characterization and Image Formation in Wide-Angle Synthetic Aperture Radar. Kush R. Varshney. *Master's Thesis*. Massachusetts Institute of Technology, Cambridge, MA, May 2006.

Prefaces

AI Ethics. Sameep Mehta, Francesca Rossi, and Kush R. Varshney, *IBM Journal of Research and Development*, vol. 63, no. 4/5, July/September 2019.

Advances in Computational Creativity Technology. Richard Goodwin, Kush R. Varshney, and Jinjun Xiong. *IBM Journal of Research and Development*, vol. 63, no. 1, January–February 2019.

Data Science for Social Good. Aleksandra Mojsilović and Kush R. Varshney. *IBM Journal of Research and Development*, vol. 61, no. 6, November–December 2017.

Online Contributions

Guest, AI for Good and the Real World, Season 4, Episode 19, *Talking Machines* (https://www.thetalkingmachines.com/episodes/ai-good-and-real-world), 18 October 2018.

Introducing AI Fairness 360. Kush R. Varshney. *IBM Research Blog* (https://www.ibm.com/blogs/research/2018/09/ai-fairness-360), 19 September 2018.

AI for Code Encourages Collaborative, Open Scientific Discovery. Kush R. Varshney. *IBM Research Blog* (https://www.ibm.com/blogs/research/2018/08/ai-code), 15 August 2018.

New AI Algorithm Recommends Right Products at the Right Time. Kush R. Varshney. *IBM Research Blog* (https://www.ibm.com/blogs/research/2017/12/ai-algorithm-recommends-products), 8 December 2017.

Reducing Discrimination in AI with New Methodology. Kush R. Varshney. *IBM Research Blog* (https://www.ibm.com/blogs/research/2017/12/ai-reducing-discrimination), 6 December 2017.

Social Good, Meet Data Science. Kush R. Varshney. *IBM Big Data & Analytics Hub Blog* (http://www.ibmbigdatahub.com/blog/social-good-meet-data-science), 16 February 2016.

For Analytics to Have an Impact, Keep it Simple. Tyler H. McCormick, Cynthia Rudin, Dmitry Malioutov, and Kush Varshney. *Data Informed* (http://data-informed.com/for-analytics-to-have-an-impact-keep-it-simple), 3 August 2015.

Statistical Analysis of Serving and Breaking First in Tennis. Kush Varshney. *Jon Wertheim's Tennis Mailbag* (http://sportsillustrated.cnn.com/tennis/news/20131030/serena-williams-mailbag), 30 October 2013.

Patents:

Issued
Method, System and Computer Program Product for Automating Expertise Management Using Social and Enterprise Data. John H. Bauer, Dongping Fang, Aleksandra Mojsilović, Karthikeyan Natesan Ramamurthy, Kush R. Varshney, and Jun Wang. US 10,643,140, May 5, 2020.

Generating Semantic Flow Graphs Representing Computer Programs. Ioana M. Baldini Soares, Aleksandra Mojsilović, Evan J. Paterson, and Kush R. Varshney. US 10,628,282, April 21, 2020.

Generating Work Products Using Work Product Metrics and Predicted Constituent Availability. Debarun Bhattacharjya, Kush R. Varshney, and Lav R. Varshney. US 10,467,638, November 5, 2019.

Accelerating Data-Driven Scientific Discovery. Flavio P. Calmon and Kush R. Varshney. US 10,388,039, August 20, 2019.

Association-Based Product Design. Kush R. Varshney, Lav R. Varshney, and Jun Wang. US 10,019,689, July 10, 2018.

Computing Personalized Probabilistic Familiarity Based on Known Artifact Data. Florian Pinel, Nan Shao, Kush R. Varshney, and Lav R. Varshney. US 9,852,380, December 26, 2017.

Active Odor Cancellation. Kush R. Varshney and Lav R. Varshney. US 9,600,793, March 21, 2017.

Food Steganography. Kush R. Varshney and Lav R. Varshney. US 9,417,221, August 16, 2016.

Predictive and Descriptive Analysis on Relations Graphs with Heterogeneous Entities. Aleksandra Mojsilović, Kush R. Varshney, and Jun Wang. US 9,195,941, November 24, 2015.

Applied
Generation and Management of an Artificial Intelligence (AI) Model Documentation Throughout Its Life Cycle. Matthew R. Arnold, Rachel K. E. Bellamy, Kaoutar El Maghraoui, Michael Hind, Stephanie Houde, Kalapriya Kannan, Manish Kesarwani, Sameep Mehta, Aleksandra Mojsilović, Ramya Raghavendra, Darrell C. Reimer, John T. Richards, David J. Piorkowski, Jason Tsay, and Kush R. Varshney.

Artificial Intelligence Certification of Factsheets Using Blockchain. Samuel C. Hoffman, Kalapriya Kannan, Pranay K. Lohia, Sameep Mehta, and Kush R. Varshney. 16/795,455, filed February 19, 2020.

Mitigating Statistical Bias in Artificial Intelligence Models. Rachel K. E. Bellamy, Kush R. Varshney, and Yunfeng Zhang. 16/745,872, filed January 17, 2020.

Enhancing Fairness in Transfer Learning for Machine Learning Models with Missing Protected Attributes in Source or Target Domains. Supriyo Chakraborty, Amanda Coston, Ziarah Mustahsan, Karthikeyan Natesan Ramamurthy, Skyler Speakman, Kush R. Varshney, and Dennis Wei. 16/692,974, filed November 22, 2019.

Semantic Queries Based on Semantic Representation of Programs and Data Source Ontologies. Ioana M. Baldini Soares, Aleksandra Mojsilović, Evan J. Paterson, and Kush R. Varshney. 16/244,447, filed January 10, 2019.

System and Method for Post Hoc Improvement of Instance-Level and Group-Level Prediction Metrics. Manish Bhide, Pranay Lohia, Karthikeyan Natesan Ramamurthy, Ruchir Puri, Diptikalyan Saha, and Kush R. Varshney. 16/214,703, filed December 10, 2018.

Distributed Platform for Computation and Trusted Validation. Nelson Kibichii Bore, Michael Hind, Eleftheria K. Pissadaki, Ravi Kiran Raman, Sekou Lionel Remy, Roman Vaculin, and Kush R. Varshney. 16/135,438, filed September 19, 2018.

Distributed Platform for Computation and Trusted Validation. Nelson Kibichii Bore, Michael Hind, Eleftheria K. Pissadaki, Ravi Kiran Raman, Sekou Lionel Remy, Roman Vaculin, and Kush R. Varshney. 16/135,326, filed September 19, 2018.

Distributed Platform for Computation and Trusted Validation. Nelson Kibichii Bore, Michael Hind, Eleftheria K. Pissadaki, Ravi Kiran Raman, Sekou Lionel Remy, Roman Vaculin, and Kush R. Varshney. 16/135,260, filed September 19, 2018.

Tracking the Evolution of Topic Rankings from Contextual Data. Mary E. Helander, Hemank Lamba, Nizar Lethif, Joana S. B. T. Maria, Emily A. Ray, and Kush R. Varshney, 15/995,653, filed May 31, 2018.

Node Relevance Determination in an Evolving Network. Mary E. Helander, Hemank Lamba, Nizar Lethif, Joana S. B. T. Maria, Emily A. Ray, and Kush R. Varshney, 15/987,048, filed May 23, 2018.

Recommendation Prediction Based on Preference Elicitation. Aleksandra Mojsilović, Kush R. Varshney, Jun Wang, and Jinfeng Yi. 15/849,214, filed December 20, 2017.

Humanitarian Crisis Analysis Using Secondary Information Gathered by a Focused Web Crawler. Ioana M. Baldini Soares, Amit Dhurandhar, Abhishek Kumar, Aleksandra Mojsilović, Kien T. Pham, Kush R. Varshney, and Maja Vukovic. 15/484,325, filed April 11, 2017.

Representation of a Data Analysis Using a Flow Graph. Ioana M. Baldini Soares, Aleksandra Mojsilović, Evan J. Paterson, and Kush R. Varshney. 15/399,420, filed January 5, 2017.

Method for Market Risk Assessment for Healthcare Applications. Shilpa Mahatma, Aleksandra Mojsilović, Karthikeyan Natesan Ramamurthy, Kush R. Varshney, Dennis Wei, and Gigi Yuen-Reed. 14/699,482, filed April 29, 2015.

Nonparametric Tracking and Forecasting of Multivariate Data. Aleksandr Y. Aravkin, Dmitry M. Malioutov, and Kush R. Varshney. 14/480,704, filed September 9, 2014, abandoned July 24, 2018.

Tutorials, Talks, Panels, and Demonstrations:

Tutorials
AI Fairness 360.
- Immersive AI Session, Open Data Science Conference, New York, NY, June 27, 2019.
- O'Reilly Artificial Intelligence Conference, New York, NY, April 16, 2019.

- ACM Conference on Fairness, Accountability, and Transparency, Atlanta, GA, January 29, 2019.

Artificial Intelligence and Machine Learning.
- United Nations Data Innovation Lab, Nairobi, Kenya, March 13, 2017.

Introduction to Business Analytics.
- IEEE International Conference on Acoustics, Speech, and Signal Processing, Kyoto, Japan, March 25, 2012.

Invited Talks
Results in Trustworthy Machine Learning That Go Against Conventional Wisdom.
- Computation and Society Initiative Colloquium, Yale University, New Haven, CT, February 21, 2020.

Computational Intelligence, Creativity, and Wisdom for Man and Mankind.
- IDInsight, Nairobi, Kenya, October 30, 2019.
- New England Complex Systems Institute Winter School, Cambridge, MA, January 11, 2018.

Developing Trust in Artificial Intelligence and Machine Learning for High-Stakes Applications.
- Mind Food Talk, Ajua (formerly mSurvey), Nairobi, Kenya, September 27, 2019.
- Carnegie Mellon University Africa, Kigali, Rwanda, September 19, 2019.
- *Keynote Address*, IEEE Syracuse Section Fellows Night, Syracuse, NY, April 11, 2019.
- Electrical Engineering and Computer Science Seminar, Syracuse University, Syracuse, NY, April 10, 2019.
- Princeton University, Princeton, NJ, March 8, 2019.
- Electrical and Computer Engineering Seminar, University of Virginia, Charlottesville, VA, October 19, 2018.
- New York Artificial Intelligence Meetup, New York, NY, September 27, 2018.

AI Fairness 360.
- Reinforce AI Conference, Budapest, Hungary, March 22, 2019.
- Budapest Data Science Meetup, Budapest, Hungary, March 20, 2019.
- Cognitive Systems Institute Group Speaker Series Webinar, September 20, 2018.

Unequal Performance Across Groups in Face Image Classification.
- SAMSI Discussion Series on Controversial Topics in Machine Learning and Precision Medicine, Research Triangle Park, NC, February 27, 2019.

AI Safety in Cyber-Physical Systems, Decision Sciences, and Data Products.
- QuantPort, New York, NY, September 13, 2018.
- Inside the Black Box Conference, New York, NY, June 16, 2018.
- Uehiro-Carnegie-Oxford Ethics Conference, New York, NY, May 17, 2018.

On Making Machine Learning Safe.
- PDT Partners, New York, NY, June 8, 2017.
- IEEE International Conference on Signal Processing and Integrated Networks, Noida, India, February 3, 2017.
- Computer Science Colloquium, University at Albany, Albany, NY, November 7, 2016.
- *Keynote Lecture*, INFORMS Southeast Michigan Symposium, Rochester, MI, May 14, 2016.

Data for Good.

- School of Social Policy and Practice, University of Pennsylvania, Philadelphia, PA, March 30, 2017.

Learning Classification Rules via Boolean Compressed Sensing with Application to Workforce Analytics.

- Optimization Seminar, Department of Operations Research and Financial Engineering, Princeton University, Princeton, NJ, January 22, 2015.
- Communications Seminar, Coordinated Science Laboratory, University of Illinois at Urbana-Champaign, Urbana, IL, November 3, 2014.
- Interdisciplinary Distinguished Seminar Series, Department of Electrical and Computer Engineering, North Carolina State University, Raleigh, NC, October 4, 2013.

Detecting Poverty with Satellite Imagery.

- Nielsen Machine Learning Webinar, December 16, 2016.
- *Keynote Address*, Big Data Summit, University of Illinois Research Park, Champaign, IL, November 5, 2014.
- Data Analysts for Social Good Webinar, October 9, 2014.

Proactive Retention among IBM Growth Market Employees.

- Using Analytics to Optimize Your Workforce Seminar, Analytics Solutions Center, Washington, DC, March 12, 2013.

Margin-Based Classification and Dimensionality Reduction Using Geometric Level Sets.

- École Centrale Paris, Châtenay-Malabry, France, July 1, 2011.
- General Electric Global Research Center, Niskayuna, NY, January 21, 2010.
- MITRE Corporation, Bedford, MA, January 14, 2010.
- NASA Jet Propulsion Laboratory, Pasadena, CA, December 10, 2009.
- MIT Lincoln Laboratory, Lexington, MA, November 24, 2009.
- IBM Thomas J. Watson Research Center, Yorktown Heights, NY, November 16, 2009.
- Naval Research Laboratory, Washington, DC, November 10, 2009.
- National Institute of Standards and Technology, Gaithersburg, MD, November 9, 2009.
- Information, Systems and Networks Seminar, Cornell University, Ithaca, NY, November 6, 2009.
- Schlumberger-Doll Research Center, Cambridge, MA, November 2, 2009.

Supervised Classification in Sensor Networks.

- Communications, Networking, Signal and Image Processing Seminar, Purdue University, West Lafayette, IN, October 6, 2010.
- Business Analytics and Mathematical Sciences Department Seminar, IBM Thomas J. Watson Research Center, Yorktown Heights, NY, October 4, 2010.

Frugal Hypothesis Testing and Classification.

- Massachusetts Institute of Technology, Cambridge, MA, February 8, 2010.

Level Set Margin-Based Classification.

- National Security Engineering Division Seminar, Lawrence Livermore National Laboratory, Livermore, CA, August 11, 2009.
- Decision Modeling Research Initiative Seminar, MIT Lincoln Laboratory, Lexington, MA, January 16, 2009.

- Johns Hopkins University Applied Physics Laboratory, Laurel, MD, January 13, 2009.

Wide-Angle SAR Image Formation with Sparsifying Regularization.
- Decision Modeling Research Initiative Seminar, MIT Lincoln Laboratory, Lexington, MA, January 30, 2007.
- Dipartimento di Ingegneria dell'Informazione, Facoltà di Ingegneria, Università di Pisa, Italy, July 17, 2006.

Panels
Panelist, Data Privacy and Ethics, Yale Economic Development Symposium, New Haven, CT, February 21, 2020.

Moderator, AI for Social Good, Yale Economic Development Symposium, New Haven, CT, February 21, 2020.

Panelist, AAAI Artificial Intelligence Diversity, Belonging, Equity, and Inclusion Workshop, New York, NY, February 7, 2020.

Panelist, NetHope Workshop, AI for Good Global Summit, Geneva, Switzerland, May 31, 2019.

Panelist, Data and Partnerships, AI XPRIZE Team and Red Judge Summit, Geneva, Switzerland, May 31, 2019.

Panelist, 'Living with AI' Future Trends Forum, Fundación Innovación Bankinter, Madrid, Spain, May 28–30, 2019.

Panelist, Internet Policy Revolution, AI and Transparency, IBM IT Services Legal Summit, New York, NY, October 15, 2018.

Panelist, SDGs and Beyond: Leveraging Technology for Development, The Youth Assembly, New York, NY, August 12, 2018.

Moderator, Machine Learning that Matters Now, Stockholm, Sweden, July 14, 2018.

Panelist, Ethics and Fairness in Artificial Intelligence Think Tank, IBM Think, Las Vegas, NV, March 22, 2018.

Panelist, IBM Watson Health Opioid Summit, Cambridge, MA, January 17, 2018.

Panelist, Human Interpretability in Machine Learning, Sydney, Australia, August 10, 2017.

Moderator, Rapid-Fire Introduction to Data Science for Social Good Organizations and Opportunities, New York, NY, June 24, 2016.

Panelist, Stories from the Frontlines of Data Science for Good, New York, NY, October 27, 2015.

Panelist, IEEE Spectrum Forecasters Panel, 2013–2015.

Demonstrations
Automatic Generation of Factsheets for Trusted AI in a Runtime Environment.

- NeurIPS Expo, Montreal, Canada, December 2, 2018.

SellerScope: Interactive Salesforce Analytics.
- IBM Information on Demand EXPO, Las Vegas, NV, October 21–24, 2012.
- IBM Investor Briefing, Yorktown Heights, NY, May 9, 2012.
- IBM Innovation Lab, Lotusphere, Orlando, FL, January 15–19, 2012.

Teaching and Mentoring Experience:

Mentor
- Gevorg Ghalachyan, M.S. candidate, *Yerevan State University*, Trustworthy AI Pentathlon, Summer–Fall 2019.
- Peter Yefi, M.S. candidate, *Carnegie Mellon University – Africa*, Trustworthy AI Pentathlon, Summer 2019.
- Timothy Odonga, M.S. candidate, *Carnegie Mellon University – Africa*, Fairness in Dermatology, Summer 2019.
- Newton Kinyanjui, M.S. candidate, *Carnegie Mellon University – Africa,* Fairness in Dermatology, Summer 2019.
- Shubham Sharma, Ph.D. candidate, *University of Texas*, Eliciting Fairness Definitions from Policymakers, Summer 2019.
- Sanghamitra Dutta, Ph.D. candidate, *Carnegie Mellon University*, Detection Theory-Based Analysis of Accuracy, Fairness, and Explainability, Summer 2019.
- Shivashankar Subramanian, Ph.D. candidate, *University of Melbourne*, Extracting Evidence of Drug Repurposing from Scientific Papers, Summer 2019.
- Akanksha Atrey, Ph.D. candidate, *University of Massachusetts*, Causal Inference to Address the Opioid Crisis, Summer 2019.
- Vivek Gupta, Ph.D. candidate, *University of Utah*, Contrastive Explanations for Natural Language Processing, Summer 2019.
- Xiufan Yu, Ph.D. candidate, *Pennsylvania State University*, Optimal Pathways out of Poverty, Summer 2019.
- Oscar Chang, Ph.D. candidate, *Columbia University*, Creating New Antibiotics of Last Resort, Summer–Fall 2018.
- Amanda Coston, Ph.D. candidate, *Carnegie Mellon University*, Fair Financial Inclusion and Transfer Learning, Summer 2018.
- Niccolò Dalmasso, Ph.D. candidate, *Carnegie Mellon University*, Pruning Neural Networks for Multiple Objectives, Summer 2018.
- Vidya Muthukumar, Ph.D. candidate, *University of California, Berkeley*, Color Blind Neural Networks, Summer 2018.
- Chirag Nagpal, Ph.D. candidate, *Carnegie Mellon University*, Causal Inference to Address the Opioid Crisis, Summer 2018.
- Ritesh Noothigattu, Ph.D. candidate, *Carnegie Mellon University*, Winning Video Games Morally, Summer 2018.
- Ravi Kiran Raman, Ph.D. candidate, *University of Illinois*, Trusted Models and Results, Summer 2018. (Now with Analog Garage.)
- Evan Patterson, Ph.D. candidate, *Stanford University*, Open Discovery Platform for a Multiple Sclerosis Cure and Semantic Understanding of Data Science Code, Summer 2016; Summer 2017–Spring 2018.
- Tejas Dharamsi, M.S. candidate, *Columbia University*, Neurology-as-a-Service for the Developing World, Fall 2017. (Now with Trifacta.)

- Jonathan Galsurkar, M.S. candidate, *Columbia University*, Smarter Sustainable Development, Summer–Fall 2017. (Now with Capital One.)
- Anand Doshi, M.S. candidate, *University of Michigan*, Emergency Food Best Practice, Summer 2017. (Now with Plenty.)
- Aditya Garg, M.S. candidate, *Columbia University*, Demystifying Social Entrepreneurship, Summer 2017. (Now with Tesla.)
- Bernat Guillén Pegueroles, Ph.D. candidate, *Princeton University*, Cognitive Policy Advisor, Summer 2017. (Now with Google.)
- Yaoli Mao, Ph.D. candidate, *Columbia University*, User Experience of Open Discovery Platforms, Summer 2017.
- Bryce Melvin, B.S. candidate, *University of Colorado*, Changing Behaviors to Conserve Energy, Summer 2017. (Now with DreamWorks Animation.)
- Timothy NeCamp, Ph.D. candidate, *University of Michigan*, Cognitive Disease Hunter, Summer 2017.
- Minh Nguyen, Ph.D. candidate, *University of Southern California*, SimplerVoice: Overcoming Illiteracy, Summer 2017.
- Hrishikesh Rao, Ph.D. candidate, *University of Michigan*, Emergency Food Best Practice, Summer 2017.
- Jinghe Zhang, Ph.D. candidate, *University of Virginia*, Combating the Opioid Crisis, Summer 2017. (Now with Target.)
- Wenyu (Wendy) Zhang, Ph.D. candidate, *Cornell University*, Cognitive Financial Advisor for Low-Wage Workers, Summer 2017.
- Yuanshuo (David) Zhao, Ph.D. candidate, *Georgia Institute of Technology*, Ask Nature for Design Inspiration, Summer 2017. (Now with Uber.)
- Caitlin Kuhlman, Ph.D. candidate, *Worcester Polytechnic Institute*, How to Foster Innovation, Summer 2016.
- Hemank Lamba, Ph.D. candidate, *Carnegie Mellon University*, Understanding the Ecospace of Philanthropic Projects, Summer 2016.
- Subhabrata Majumdar, Ph.D. candidate, *University of Minnesota*, Hunting Zika with Machine Learning, Summer 2016. (Now with AT&T Labs Research.)
- Kien Pham, Ph.D. candidate, *New York University*, Real-Time Understanding of Humanitarian Crises, Summer 2016. (Now with Facebook.)
- Yumeng Tao, Ph.D. candidate, *University of California Irvine*, Disseminating the Best Treatment for Diarrhea, Summer 2016. (Now with Facebook.)
- Guolong Su, Ph.D. candidate, *Massachusetts Institute of Technology*, Interpretable Two-Level Rule Learning, Summer 2015. (Now with Google.)
- Michael Wattendorf, B.S. candidate, *Princeton University*, Predicting Outcomes of Individual Tennis Points, Spring 2015–Spring 2016. (Now with Waymo.)
- Ben Quazzo, B.S. candidate, *Princeton University*, Statistical Analysis of Professional Tennis, Spring 2015. (Now with Accel Partners.)
- Shaobo Han, Ph.D. candidate, *Duke University,* Nonparametric Bayesian Modeling of Legislation for Exploratory Analysis, Summer 2014. (Now with NEC Laboratories America.)
- Jinfeng Yi, Ph.D. candidate, *Michigan State University*, Expertise Assessment Recommendation, Summer 2013. (Now with JD AI Research.)
- Alex Gittens, Ph.D. candidate, *California Institute of Technology*, Voluntary Attrition Modeling, Summer 2012. (Now with Rensselaer Polytechnic Institute.)
- Gautam K. Bhat, M.Sc. candidate, *Karnataka State Open University*, Deconvolving the Productivity of Salespeople, Fall 2011–Summer 2012. (Now with IBM Global Technology Services.)

- John Z. Sun, Ph.D. candidate, *Massachusetts Institute of Technology*, Dynamic Matrix Factorization and Collaborative Filtering based on the Kalman Filter, Summer 2011. (Now with PDT Partners.)
- Ankan Saha, Ph.D. candidate, *University of Chicago*, Semi-Supervised Multi-Task Learning with Task-Dependent Regularization, Summer 2010. (Now with LinkedIn.)

Adjunct project advisor, DS-GA 1003 Machine Learning and Computational Statistics, *New York University*, Spring 2015, Spring 2016.

Guest lecturer
- PADM 5472 Leveraging Information Technology for Public and Nonprofit Management, *Cornell University*, Spring 2019, Spring 2020.
- INFO-UB 0001 Information Technology in Business and Society, *New York University*, Spring 2019.
- CS-GY 9223 Foundations of Data Science, *New York University*, Spring 2016.
- 01:090:101 Data: What is it Good For? (Absolutely Something), *Rutgers University*, Fall 2015.
- ORF 360 Decision Modeling in Business Analytics, *Princeton University*, Spring 2015.

Teaching assistant, 6.972 Algorithms for Estimation and Inference, *Massachusetts Institute of Technology*, Fall 2006.

Group tutor, ECE 310 Introduction to Probability and Random Phenomena, *Cornell University*, Spring 2004.

Guest instructor, How Television Works, *Tompkins Seneca Tioga BOCES Community (alternative) School*, Spring 2004.

Service and Professional Activities:

Founding co-director: IBM Science for Social Good initiative, since 2015.

Board of directors: CriticaLink, since 2016.

Data ambassador: DataKind, Brooklyn, NY, USA, since August 2013.

Advisor: ALCF Concept to Clinic Lung Cancer Early Detection Challenge Hosted by DrivenData, 2017.

Professional societies: IEEE, since 2000 (senior member since 2015); ACM, 2013–2016; INFORMS, 2012; ISIF, 2009–2011; SIAM, 2006–2010; SPIE, 2006–2007; IEE/IET, 2005–2006.

Elected professional committees and working groups: Partnership on AI Working Group on Safety-Critical AI, since 2018; IEEE Signal Processing Theory and Methods Technical Committee, 2015–2017; IEEE Machine Learning for Signal Processing Technical Committee, 2014–2016.

Journal editorial board: Digital Signal Processing: A Review Journal, 2013–2016.

Proposal review panel: Directorate for Computer and Information Science and Engineering, National Science Foundation, 2019.

Workshop and symposium organizer: 2016–2018, 2020 ICML Workshop on Human Interpretability (WHI) in Machine Learning; 2020 IEEE CVPR Workshop on Fair, Data Efficient and Trusted Computer Vision; 2019 ICML Workshop on Human in the Loop Learning (HILL); 2019 AAAS Annual Meeting Symposium on Blockchain and the Scientific Method; 2017 IEEE GlobalSIP Symposium on Signal and Information Processing for Finance and Business; 2017 ACM SIGKDD Workshop on Machine Learning for Creativity; 2016 ICML Workshop on #data4good: Machine Learning in Social Good Applications.

Track chair: ACM Conference on Fairness, Accountability and Transparency (FAT*), 2020.

Journal reviewer: Computer Methods in Biomechanics and Biomedical Engineering; Computers in Human Behavior; Digital Signal Processing (2 articles); Electronic Journal of Statistics; Entropy; Future Business Journal; IBM Journal of Research and Development (2 articles); IEE Proceedings Vision, Image and Signal Processing; IEEE Geoscience and Remote Sensing Letters; IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing; IEEE Journal of Selected Topics in Signal Processing (5 articles); IEEE Potentials Magazine (several articles); IEEE Signal Processing Letters (5 articles); IEEE Signal Processing Magazine (2 articles); IEEE Transactions on Aerospace and Electronic Systems; IEEE Transactions on Communications; IEEE Transactions on Cybernetics; IEEE Transactions on Image Processing (2 articles); IEEE Transactions on Information Theory (2 articles); IEEE Transactions on Signal Processing (9 articles); Information Sciences; Inverse Problems and Imaging; Journal of Artificial Intelligence Research; Journal of Machine Learning Research (4 articles); Nature Machine Intelligence; Pattern Recognition; Remote Sensing Letters; Service Science (2 articles); Signal Processing (2 articles); Socio-Economic Planning Sciences; Stochastic Models.

Conference program committee and reviewer: ICLR Workshop on Practical ML for Developing Countries, 2020; IEEE International Conference on Acoustics, Speech, and Signal Processing, 2014–2017, 2019–2020; NeurIPS Workshop on Machine Learning for Development, 2019; KDD Workshop on Adversarial Learning Methods for Machine Learning and Data Mining, 2019; IEEE/IFIP DSN Workshop on Dependable and Secure Machine Learning, 2018–2019; ICLR Workshop on Debugging Machine Learning Models, 2019; ICLR Workshop on Safe Machine Learning: Specification, Robustness, and Assurance, 2019; International Conference on Artificial Intelligence and Statistics, 2019; ACM Conference on Fairness, Accountability, and Transparency, 2019; Conference on Neural Information Processing Systems, 2016–2018; ACM SIGCAS Conference on Computing and Sustainable Societies, 2018; ACM SIGKDD Conference on Knowledge Discovery and Data Mining, 2015–2018; International Conference on Machine Learning, 2012–2014, 2018; International Conference on Learning Representations, 2018; IEEE ICDM Workshop on Data Science for Human Capital Management, 2017; IEEE International Workshop on Machine Learning for Signal Processing, 2014–2017; IEEE International Conference on Intelligent Transportation Systems, 2017; International Conference on Signal Processing and Integrated Networks, 2015–2017; NIPS Workshop on Interpretable Machine Learning in Complex Systems, 2016; IEEE International Workshop on Statistical Signal Processing, 2012, 2016; Workshop on Compressed Sensing Theory and its Applications to Radar, Sonar and Remote Sensing, 2015; IEEE Sensor Array and Multichannel Signal Processing Workshop, 2014; IEEE International Symposium on Information Theory, 2014; IEEE International Conference on Service Operations and Logistics, and Informatics, 2013; IEEE GlobalSIP Symposium on Signal and Information Processing in Finance and Economics, 2013; European Signal Processing Conference, 2009, 2012; IEEE International Workshop on Statistical Signal Processing, 2011; International Conference on Industrial, Engineering, and Other Applications of Applied Intelligent Systems, 2011.

Conference session chair: IEEE International Workshop on Machine Learning for Signal Processing, 2016; IEEE International Workshop on Statistical Signal Processing, 2012; INFORMS Annual Meeting, 2011.

Writer: Information Ashvins Blog (http://informationashvins.wordpress.com), since 2010; LIDS-Blog (http://lidsblog.typepad.com), 2007–2010.

Technical advisor, Fighter Steel Productions LLC, 2018.

Network member, *Council for Big Data, Ethics, and Society*, January–June 2016.

Organizing committee member, *Ebola Open Data Jam*, October 2014 and February 2015.

Summer intern coordinator, *Business Analytics and Mathematical Sciences Department*, IBM Thomas J. Watson Research Center, May–October 2013.

Organizing committee chair, *Stochastic Systems Symposium in Honor of Alan Willsky on the Occasion of his Sixtieth Birthday*, May 2008.

Organizer, *Stochastic Systems Group Seminars*, Fall 2007–Spring 2008.

Executive committee member, *MIT EECS Graduate Students Association*, 2006–2007.

Conference committee member, *LIDS Student Conference*, January 2005 and 2006.

Selected Coursework:

IBM Research: Technical Leadership; MicroMBA Program I & II.

Massive Open Online Courses: Quantitative Methods in Clinical and Public Health Research (HarvardX); The Challenges of Global Poverty (MITx).

Massachusetts Institute of Technology: Detection, Estimation, and Stochastic Processes; Discrete-Time Signal Processing; Optimization Methods; Machine Learning (listener); Representation and Modeling for Image Analysis; Computational Biology: Genomes, Networks, Evolution; Language and Its Structure II: Syntax; Language and Its Structure III: Semantics and Pragmatics.

Cornell University: Introduction to Computer Graphics; Signal Representation and Modeling; Digital Signal Processing; Telecommunications Systems I & II; Introduction to Scientific Computation; Geographic Information Systems; Introduction to Scanning Electron Microscopy; Computer Organization.

References:

Available upon request.