Interpreting Classifiers through Attribute Interactions in datasets

Andreas Henelius, Kai Puolamäki, Antti Ukkonen

Finnish Institute of Occupational Health
Helsinki, Finland

ICML Workshop on Human Interpretability in Machine Learning 2017

Finnish Institute of Occupational Health

BACKGROUND

State-of-the-art classifiers are high-performing but essentially **black boxes.**

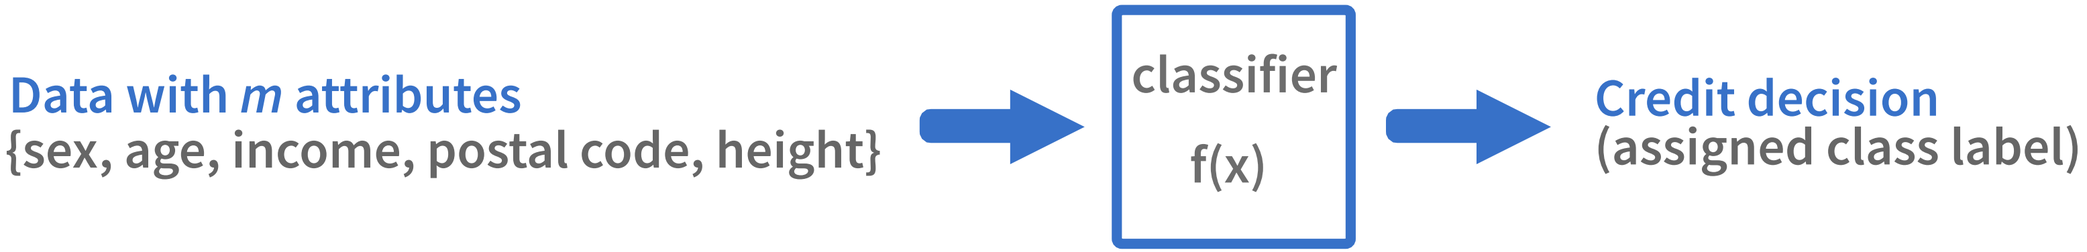

Good predictions, but why?

BACKGROUND

Need for algorithmic transparency

- decision making (e.g., possible legislative demands)
- to learn about the data (why was the class assigned?)

How can we gain insight into how the algorithms work?

INTERPRETABILITY
through attribute interactions

Interaction

Attributes are conditionally dependent given the class, i.e., attributes are together meaningful when predicting the class.

INTERPRETABILITY
through attribute interactions

Applications

bioinformatics
interactions between SNPs

pharmacogivilance
interactions between simultaneously administered drugs

INTERPRETABILITY
through attribute interactions

Grouping of attributes $\mathcal{S}$

$\mathcal{S}$ = { {sex, age}, {income, postal code}, {height}}

- attributes in the same group interact,
- attributes in different groups are independent

INTUITION OF CLASSIFIERS

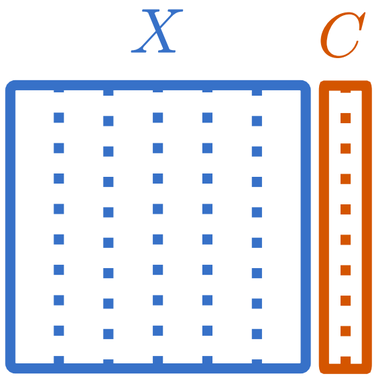

Dataset: $D = (X, C)$

Grouping of attributes: $\mathcal{S}$

The classifier models $P(C|X) \propto \underbrace{P(X|C)}_{\text{class-conditional distribution}} P(C)$

Factorise the class-conditional joint data distribution (parametrised by $\mathcal{S}$)

$$P(X \mid C; \mathcal{S}) = \prod_{S \in \mathcal{S}} P(X(\cdot, S) \mid C)$$

INTUITION OF CLASSIFIERS

Assumption: classification accuracy decreases monotonically
Known grouping: {{1, 2}, {3}, {4}}

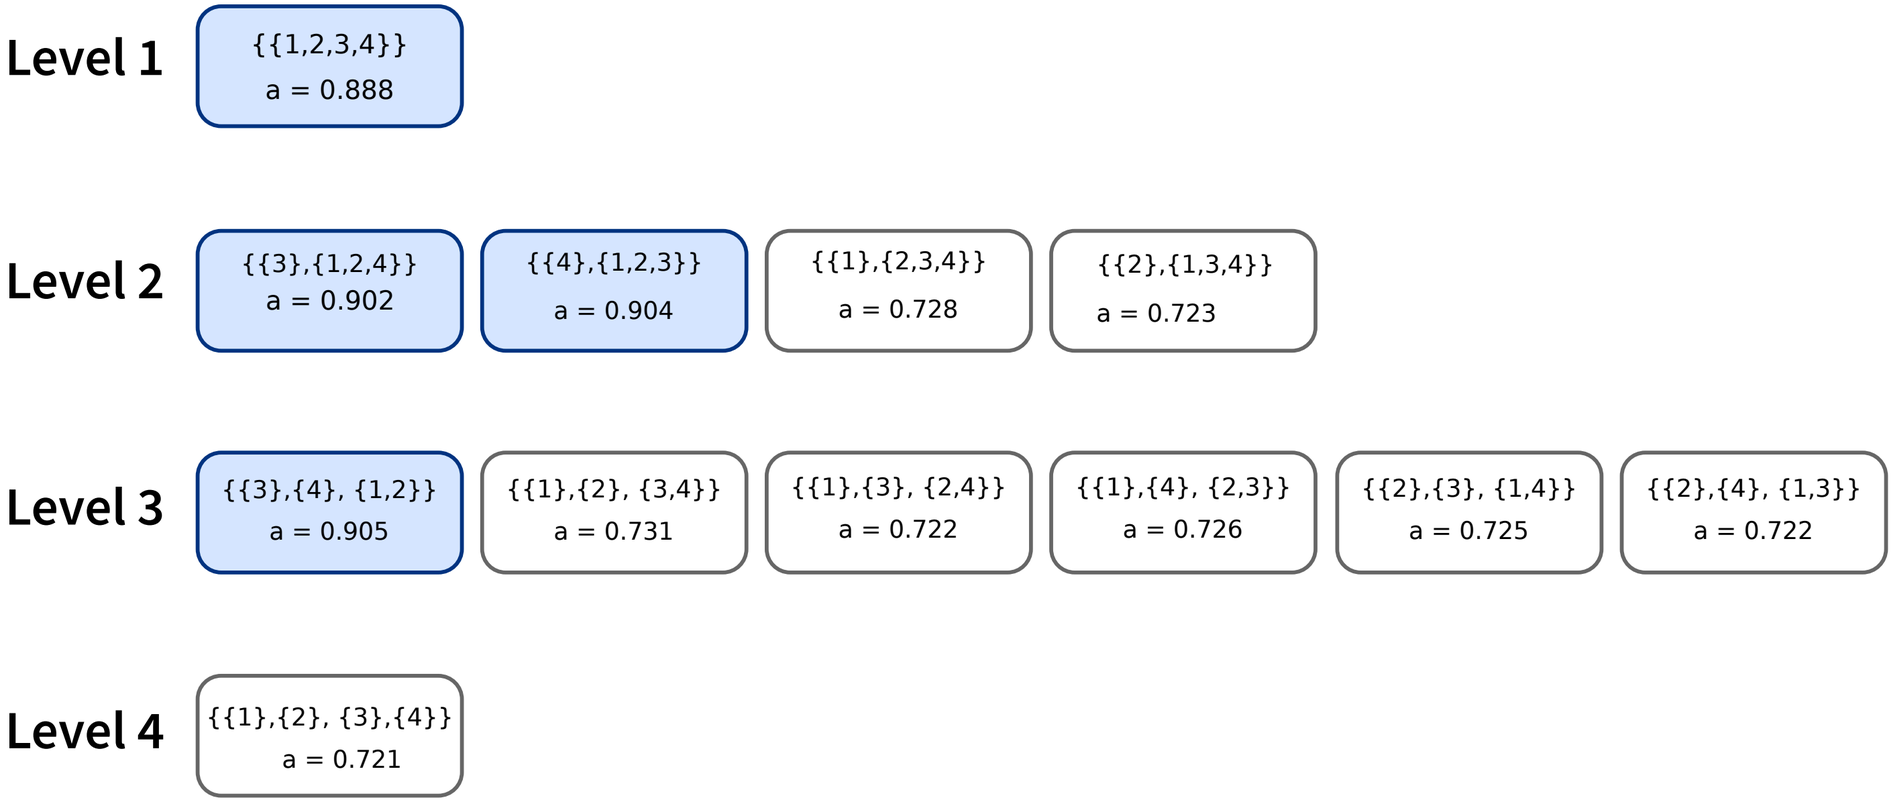

GOALS

$\mathcal{S}$ = { {sex, age}, {income, postal code}, {height}}

(1) **Verify** if a given grouping of attributes represents the attribute interaction structure

(2) **Automatically find** the maximum cardinality grouping of the attributes in the dataset

PROBLEM 1: VERIFY A GIVEN GROUPING

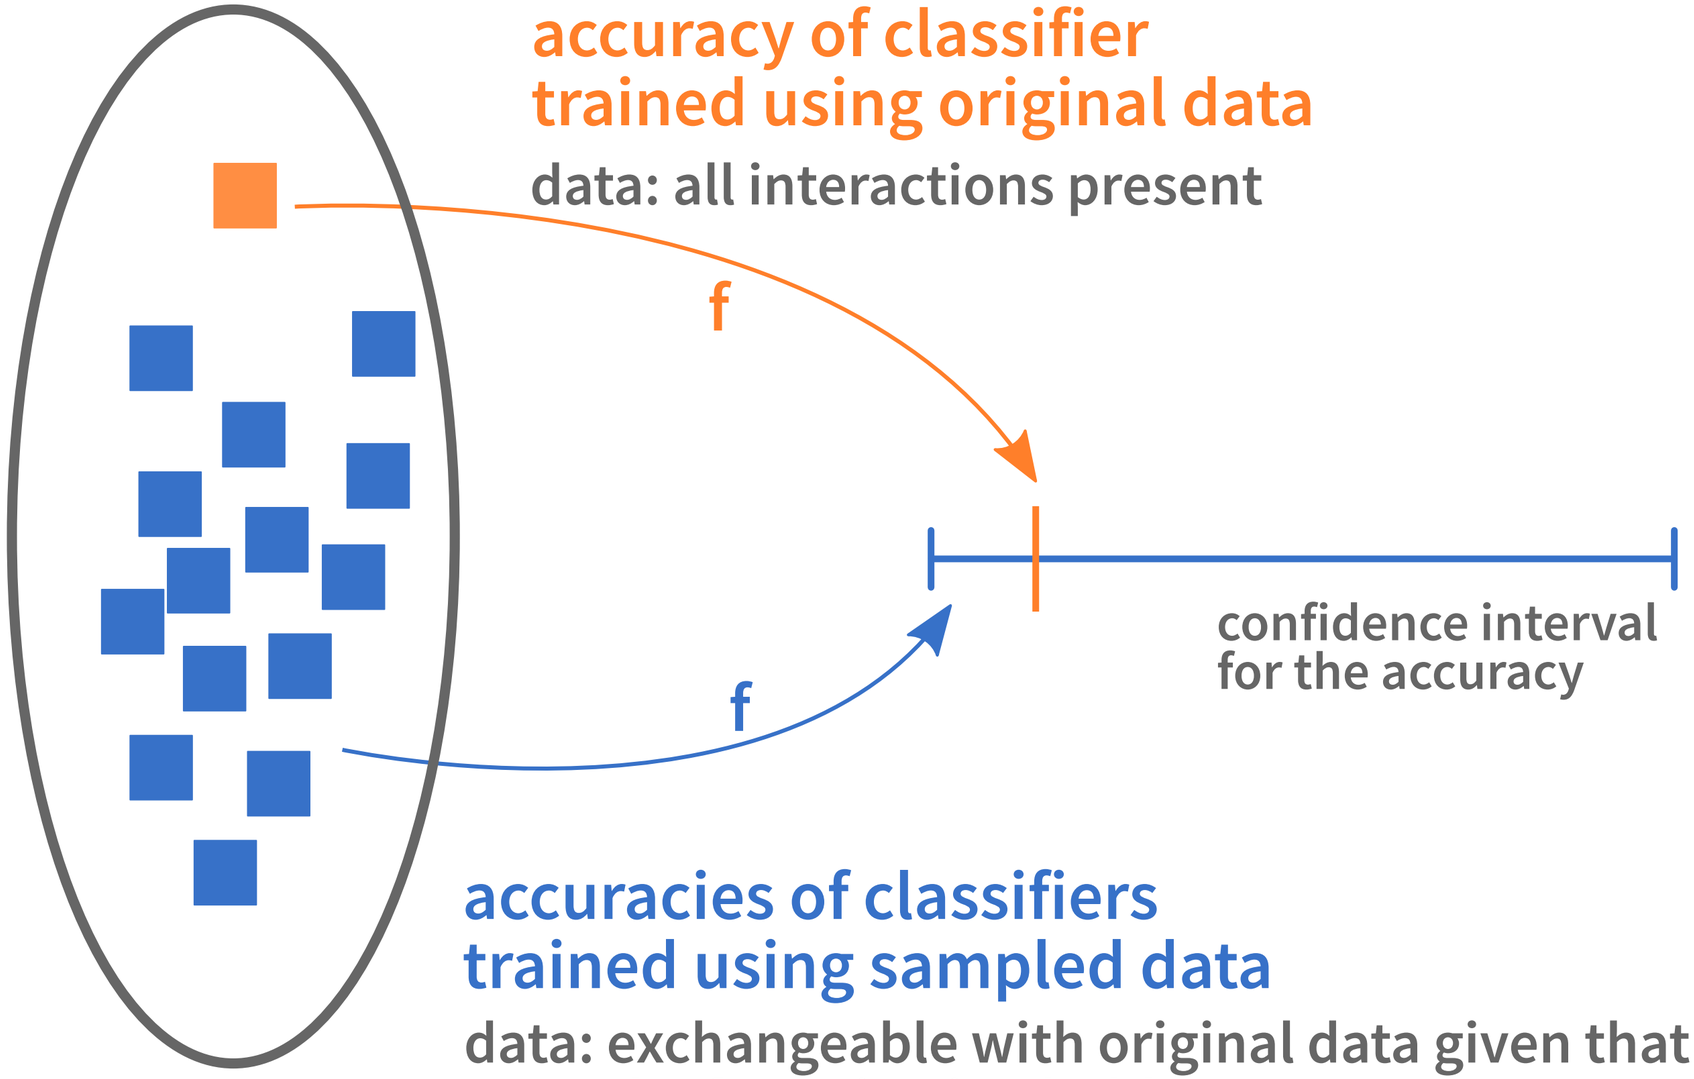

PROBLEM 1: VERIFY A GIVEN GROUPING

How do we get the data needed to construct the confidence intervals?

GENERATING EXCHANGEABLE DATASETS

Constrained randomisation, parametrised by the attribute grouping $\mathcal{S}$

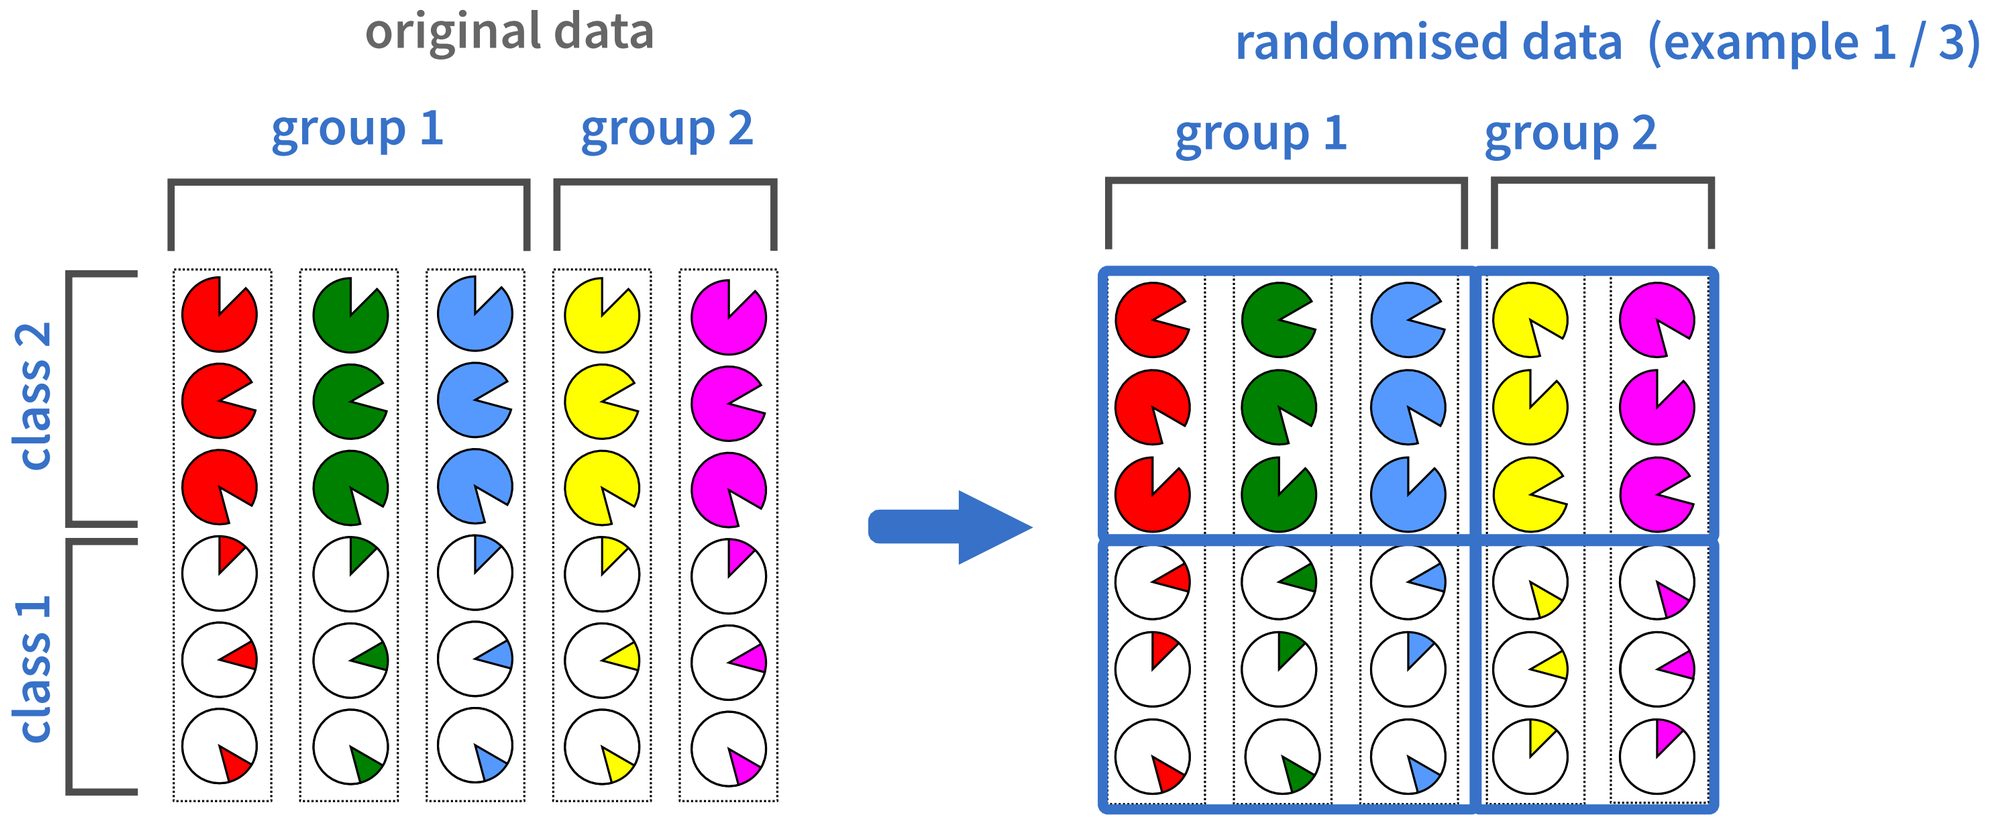

GENERATING EXCHANGEABLE DATASETS

Constrained randomisation, parametrised by the attribute grouping $\mathcal{S}$

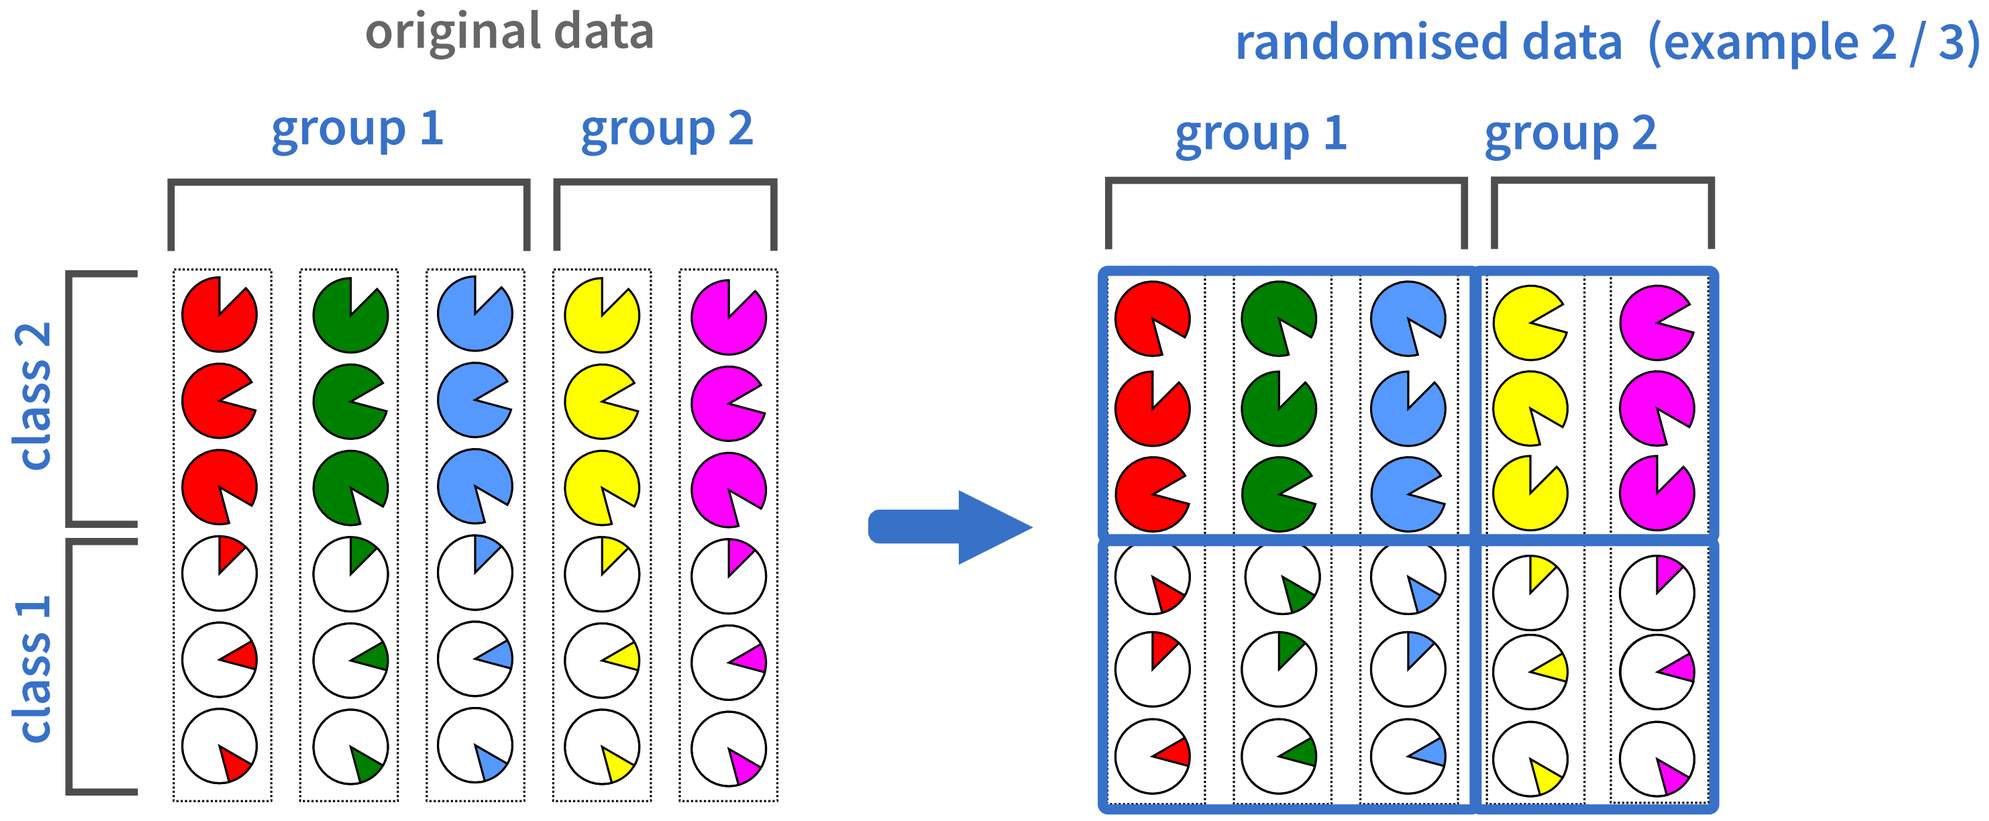

GENERATING EXCHANGEABLE DATASETS

Constrained randomisation, parametrised by the attribute grouping $\mathcal{S}$

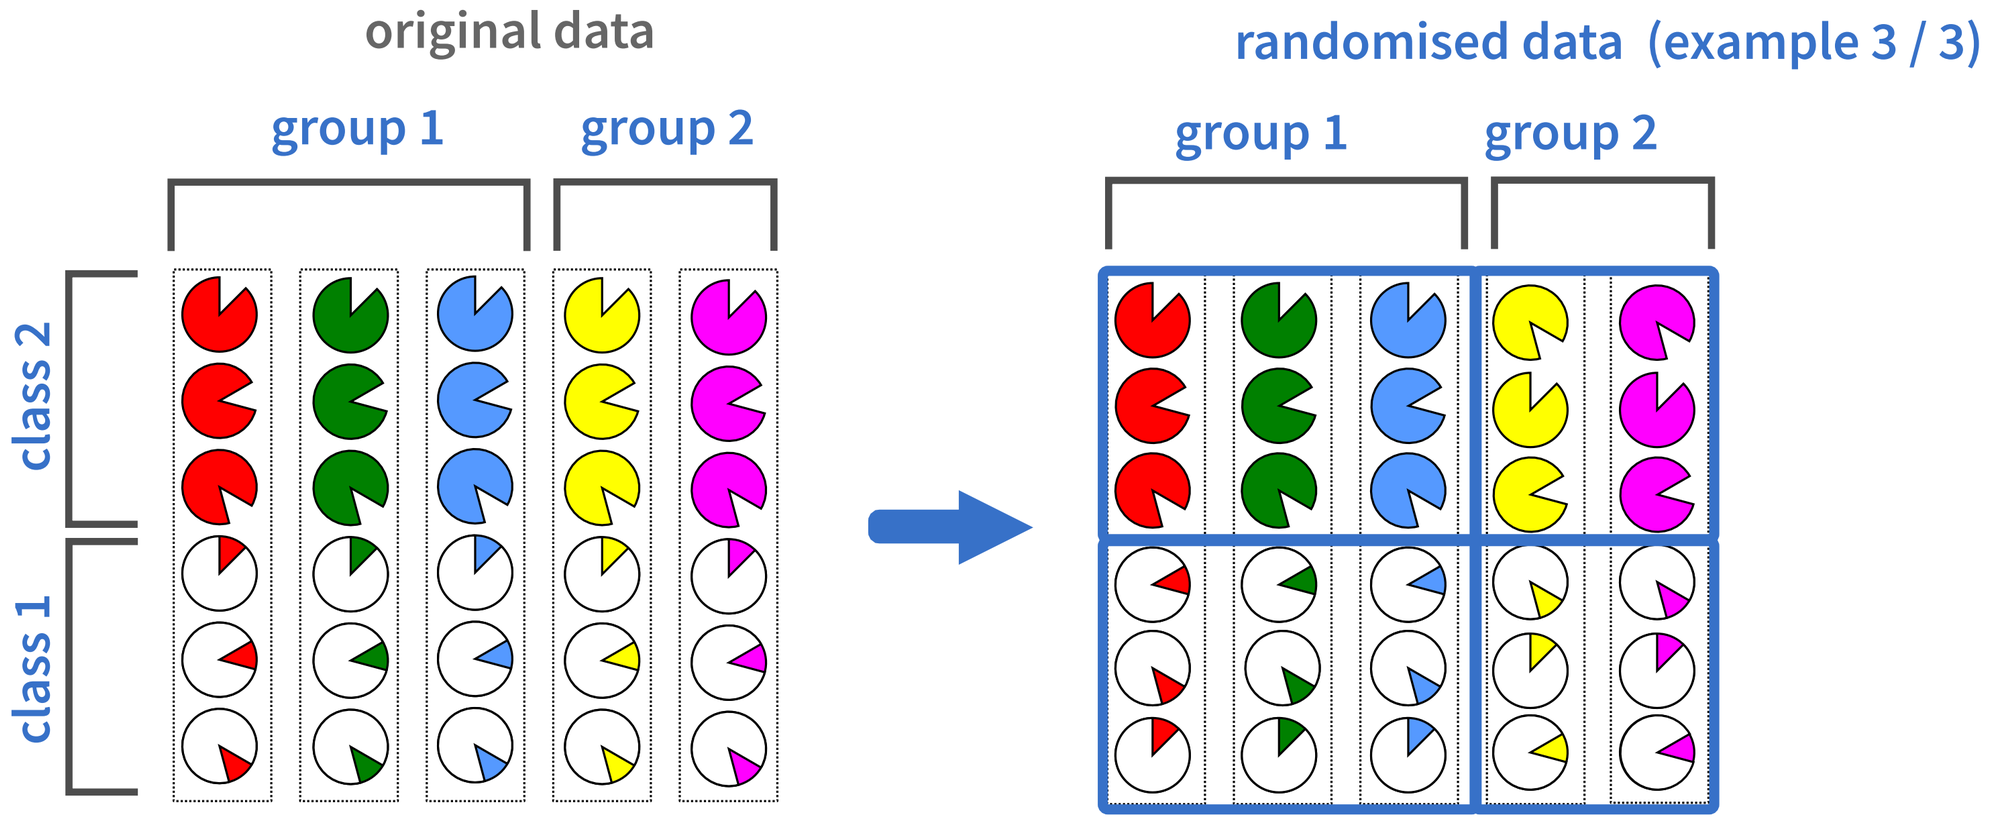

PROBLEM 2: FIND GROUPING

Top-down greedy algorithm: ASTRID
Two steps: sorting and grouping

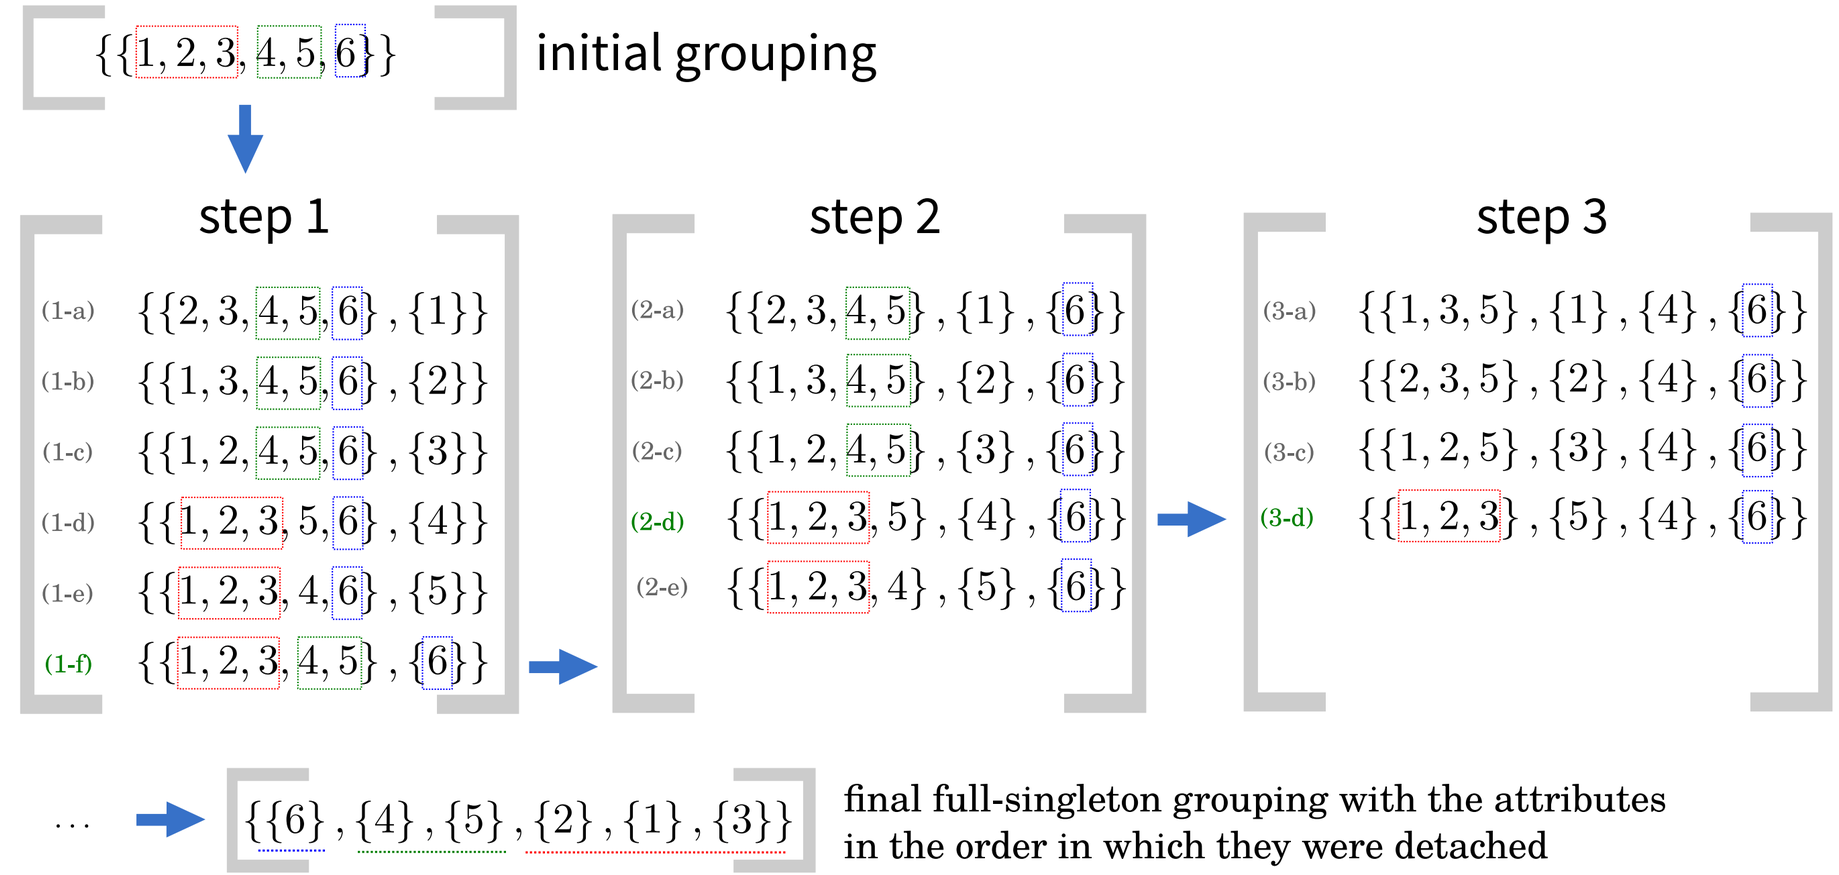

PROBLEM 2: FIND GROUPING

grouping step

$\{\{6, 4, 5, 2, 1, 3\}\}$

full-singleton grouping with the attributes in the order in which they were detached in the sorting step

(a) $\{\{6\}, \{4, 5, 2, 1, 3\}\}$ a_1

(b) $\{\{6, 4\}, \{5, 2, 1, 3\}\}$ a_3

(c) $\{\{6, 4, 5\}, \{2, 1, 3\}\}$ a_2

(d) $\{\{6, 4, 5, 2\}, \{1, 3\}\}$ a_4

(e) $\{\{6, 4, 5, 2, 1\}, \{3\}\}$ a_5

$a_1 \geq a_2 \geq \cdots \geq a_5$

→ optimal 2-grouping

a_1 $\{\{6\}, \{4, 5, 2, 1, 3\}\}$

→ optimal 3-grouping

a_1 $\{\{6\}, \{4, 5, 2, 1, 3\}\}$

a_2 $\{\{6, 4, 5\}, \{2, 1, 3\}\}$

→ $\{\{6\}, \{4, 5\}, \{2, 1, 3\}\}$

RESULTS

SVM

$a_0 = 0.908$

k	CI	a3	a4	a2	a1
2	[0.900, 0.920] *	(A)	(B	B	B)
3	[0.896, 0.920] *	(A)	(B)	(C	C)
4	[0.696, 0.784]	(A)	(B)	(C)	(D)

$\mathcal{S} = \{\{1, 2\}, \{3\}, \{4\}\}$

Random Forest

$a_0 = 0.904$

k	CI	a3	a4	a1	a2
2	[0.896, 0.928] *	(A)	(B	B	B)
3	[0.896, 0.928] *	(A)	(B)	(C	C)
4	[0.668, 0.756]	(A)	(B)	(C)	(D)

$\mathcal{S} = \{\{1, 2\}, \{3\}, \{4\}\}$

Naive Bayes

$a_0 = 0.760$

k	CI	a1	a2	a3	a4
2	[0.760, 0.760] *	(A)	(B	B	B)
3	[0.760, 0.760] *	(A)	(B)	(C	C)
4	[0.760, 0.760] *	(A)	(B)	(C)	(D)

$\mathcal{S} = \{\{1\}, \{2\}, \{3\}, \{4\}\}$

Known grouping: {{1, 2}, {3}, {4}}

SUMMARY

- Investigate attribute interactions used by classifiers
- Find groups of interacting attributes
- Enhances interpretability of opaque classifiers
- R-package: **https://github.com/bwrc/astrid-r/**

Thank you!